

Introduction to Deep Learning

ICME Summer Workshops 2021

Session 1: 8:00–9:30 AM

Instructor: Sherrie Wang

icme-workshops.github.io/deep-learning



Zakharov et al. "Few-Shot Adversarial Learning of Realistic Neural Talking Head Models" (2019)

Workshop Schedule

Day 1

Session 1 (Today 8:00–9:30 AM)

- Introduction
- Examples of deep learning
- Math review
- Neural network basics

Session 2 (Today 9:45 – 11:00 AM)

- Loss functions
- Gradient descent
- Backpropagation
- Walkthrough of an example

Day 2

Session 3 (Tomorrow 8:00–9:30 AM)

- Overfitting and underfitting
- Convolutional neural networks
- Recurrent neural networks

Session 4 (Tomorrow 9:45 – 11:00 AM)

- Other architectures
- Deep learning libraries
- Walkthrough of an example
- Failures of deep learning

Learning Goals

Day 1

Session 1 (Today 8:00–9:30 AM)

- Introduction
- Examples of deep learning
- Math review
- Neural network basics

Session 2 (Today 9:45 – 11:00 AM)

- Loss functions
- Gradient descent
- Backpropagation
- Walkthrough of an example

- What deep learning is
- Building blocks of neural networks

- How neural networks learn from data

Learning Goals

- More complex neural networks
- Limitations and pitfalls when training neural networks

Day 2

Session 3 (Tomorrow 8:00–9:30 AM)

- Overfitting and underfitting
- Convolutional neural networks
- Recurrent neural networks

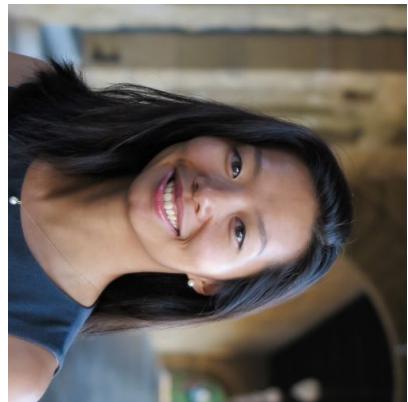
Session 4 (Tomorrow 9:45–11:00 AM)

- Other architectures
- Deep learning libraries
- Walkthrough of an example
- Failures of deep learning

Workshop Logistics

- Workshop website:
<https://icme-workshops.github.io/deep-learning/>
- Slides for today's lectures are available online
- We will be using Slido for polls throughout lecture:
<https://tinyurl.com/dlworkshop2021>
- Please post all questions during lecture to Piazza:
<https://piazza.com/icme/summer2021/icme>
- At the end of each session, I will take questions over Zoom
- New 2-day format! I encourage everyone to ask questions so that on day 2 all previous concepts are clear
- Will be recorded and uploaded online

About me



Hello
my name is

Sherrie Wang

- PhD graduate of ICME (2021)
- **Research focus:** Machine learning methods for remote sensing and applications in sustainability
- **Relevant courses:**
 - CS 221 (Artificial Intelligence)
 - CS 229 (Machine Learning)
 - CS 228 (Probabilistic Graphical Models)
 - CS 230 (Deep Learning)
 - CS 231n (Convolutional Neural Networks)
 - CS 236 (Generative Adversarial Networks)
 - CS 330 (Deep Multi-Task and Meta Learning)
- **Relevant teaching:** CME 250 (Introduction to Machine Learning), Summer Workshop 2019-20

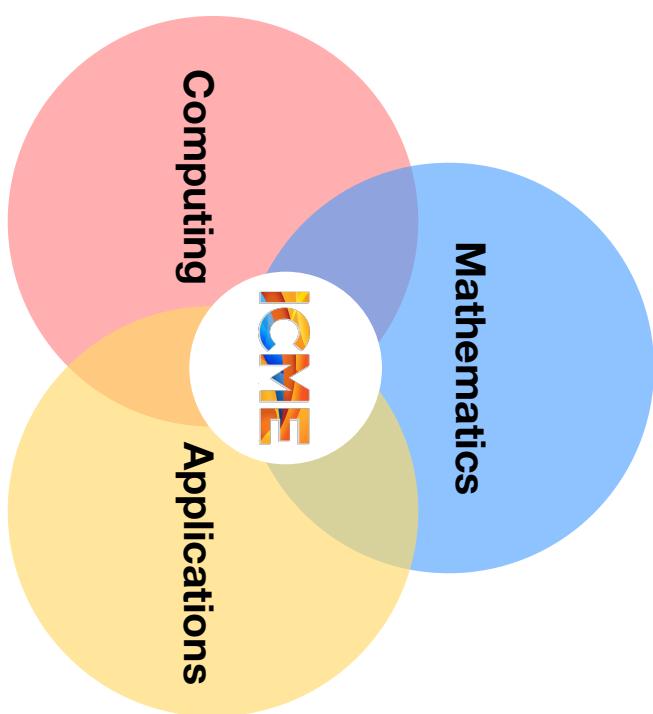


What is **ICME**?

Stanford Institute for Computational and Mathematical Engineering

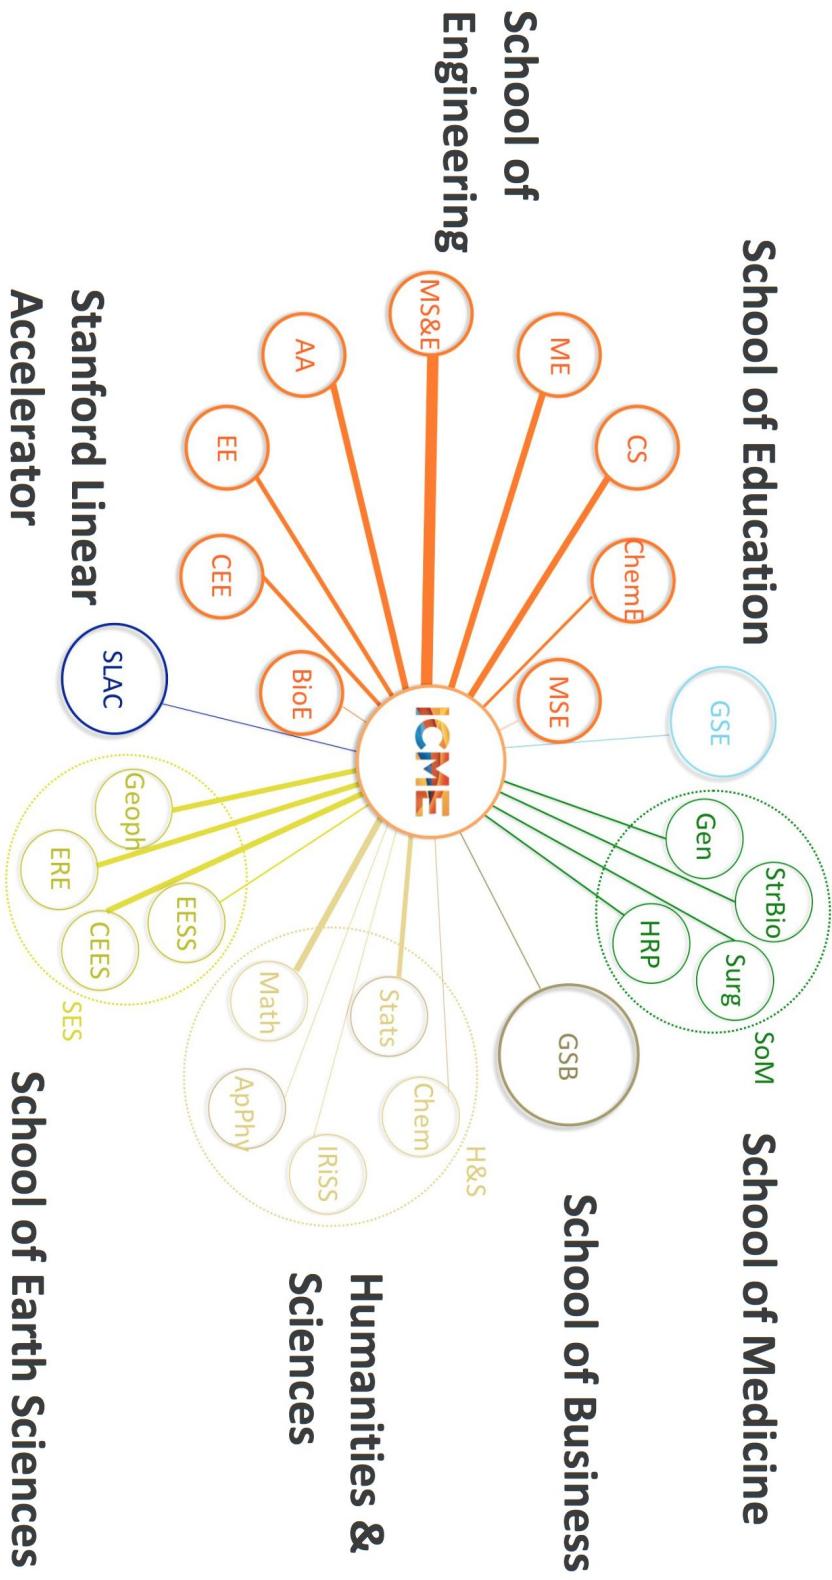
Hub at Stanford

Interdisciplinary field using
advanced mathematics
and computing to address
complex problems



icme.stanford.edu

55+ affiliated faculty across Stanford



150+ MS and PhD students



Introduction to Deep Learning - ICML Summer Workshops 2021

ICME industry-academia ecosystem



WELLS
FARGO

TOTAL

Intuit

Microsoft

NVIDIA®

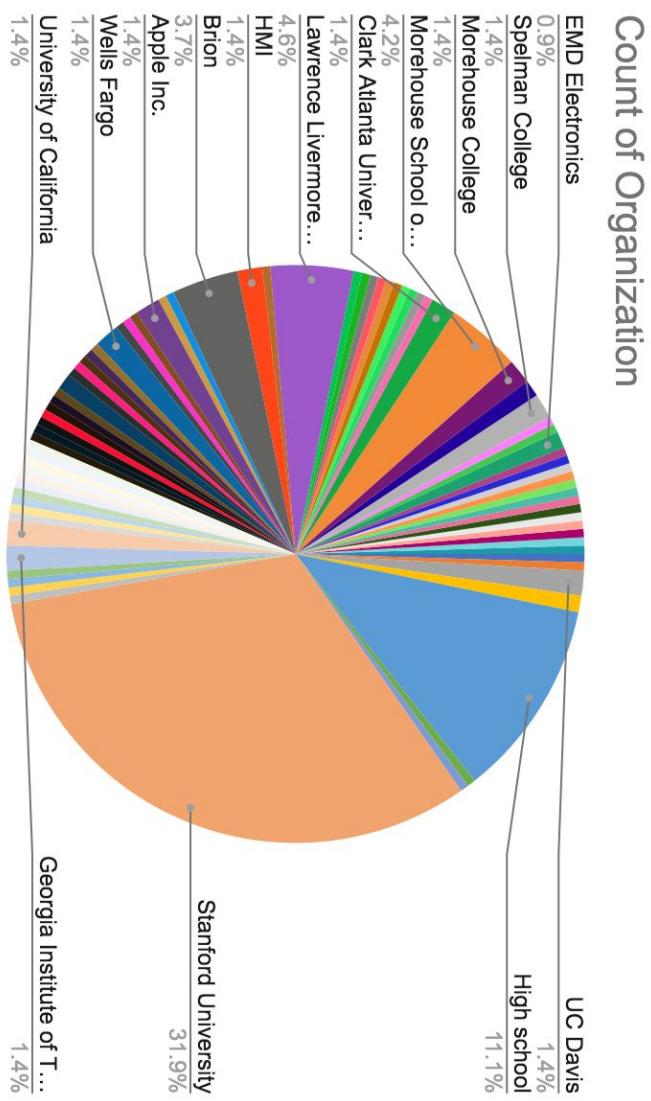
Schlumberger Tencent 腾讯 Google MathWorks

@WalmartLabs

Accenture Applied Intelligence

Who's here today

200+ registered participants



The Deep Learning Revolution

What is “deep learning”?

Artificial intelligence

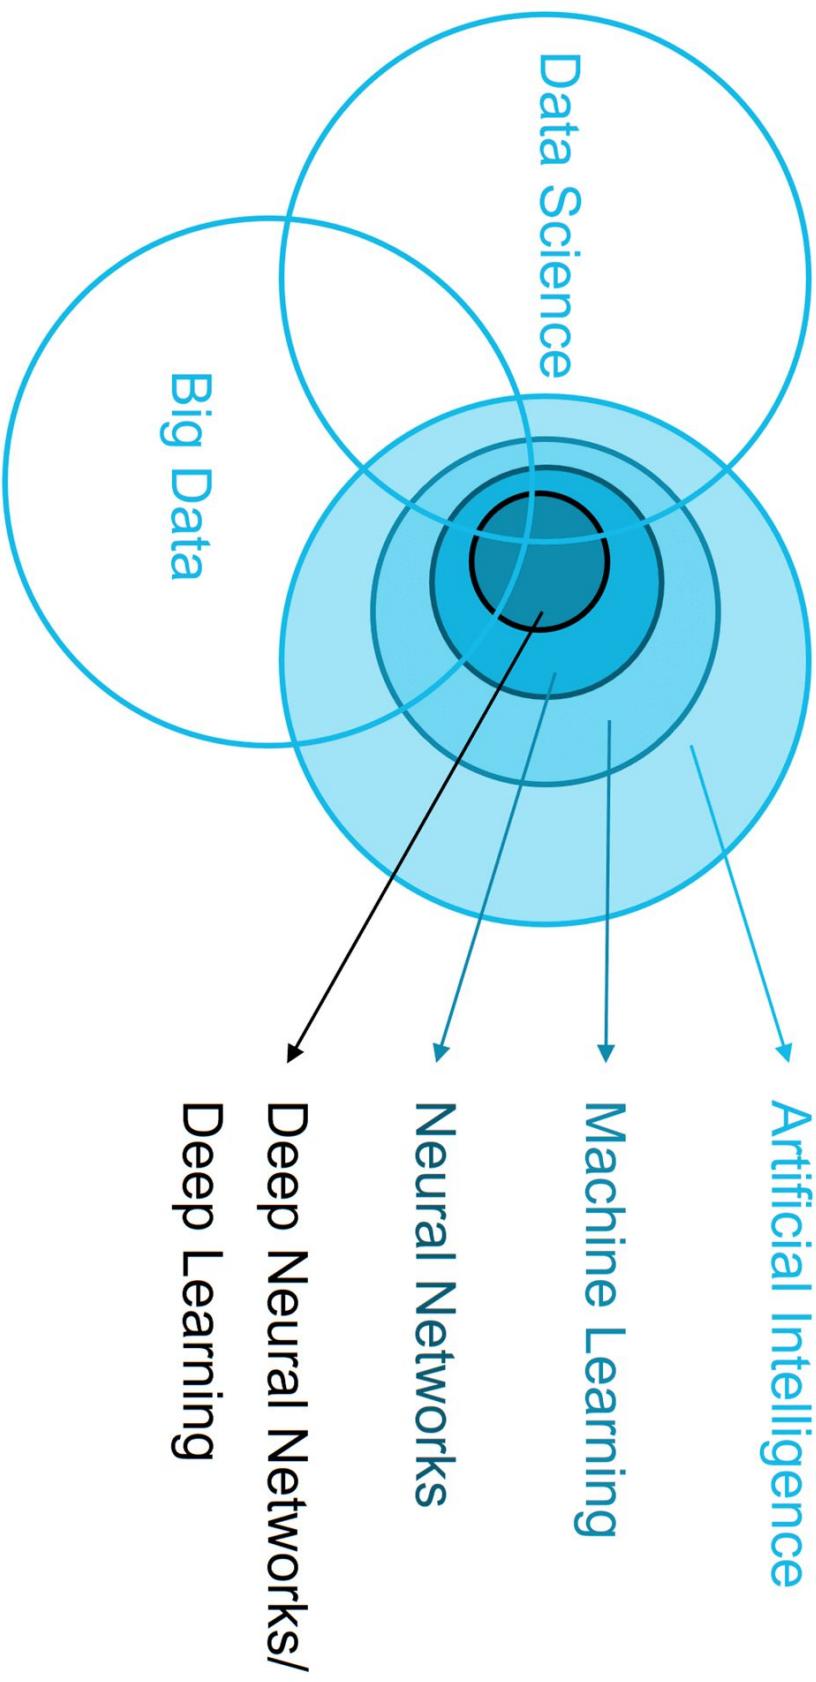
Big data

Machine learning

Deep learning

Data science

What is “deep learning”?



<https://tamuk.quickstart.com/blog/artificial-intelligence-vs.-machine-learning-vs.-data-science/>

Introduction to Deep Learning - ICME Summer Workshops 2021

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

Computer vision is one of the first places where deep learning started to outshine other methods.



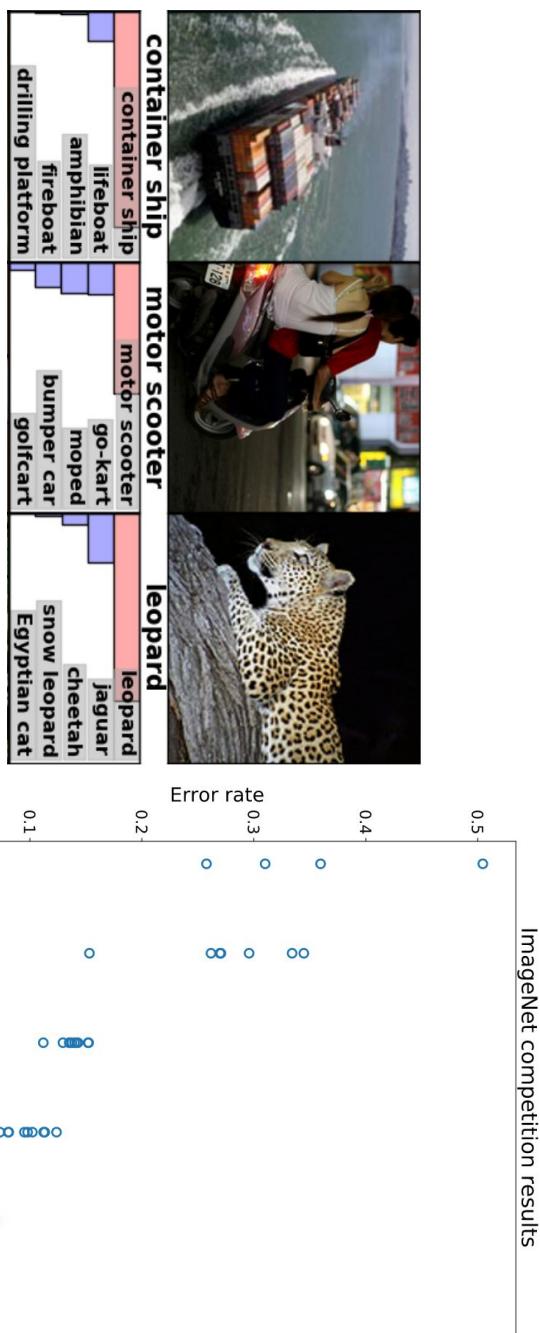
"The Street View House Numbers (SVHN) Dataset": <http://ufldl.stanford.edu/housenumbers/>
Introduction to Deep Learning - ICME Summer Workshops 2021

1. Optical character recognition

Deep learning examples

1. Optical character recognition (OCR)
2. **Image classification**
3. Text translation
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

2. Image classification



Krizhevsky, Sutskever, and Hinton. "ImageNet Classification with Deep Convolutional Neural Networks" (2012)
Introduction to Deep Learning - ICME Summer Workshops 2021

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. **Text translation**
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

3. Text translation

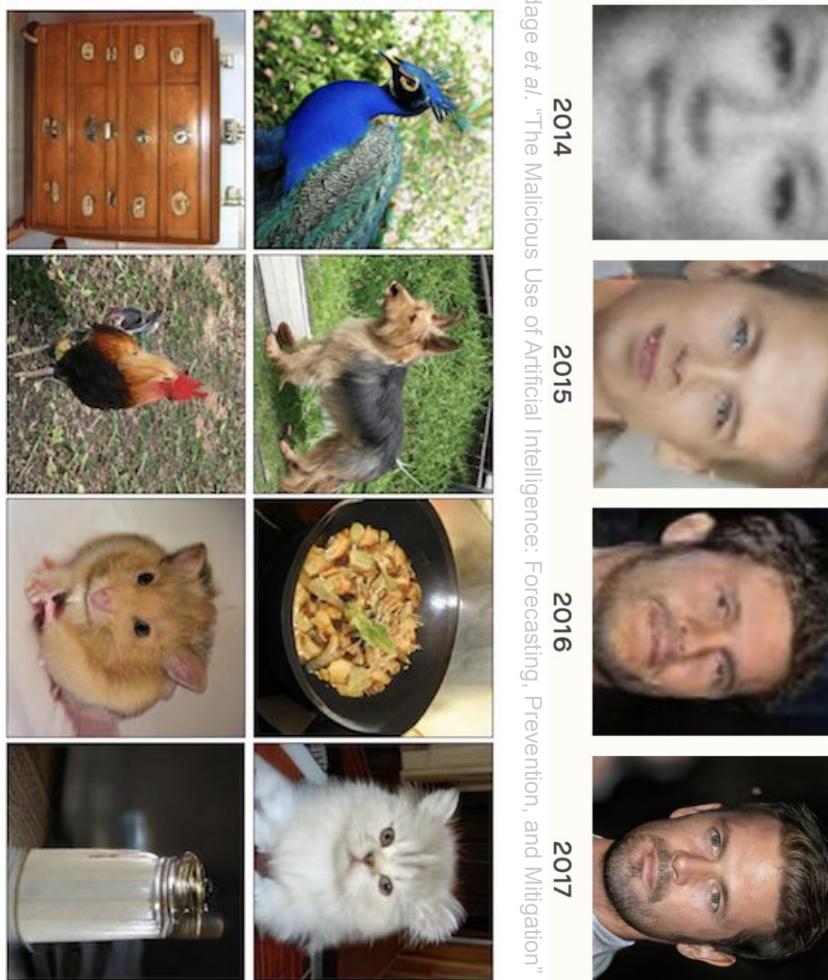


Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. **Image generation**
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

4. Image generation

Progression in the capabilities of GANs



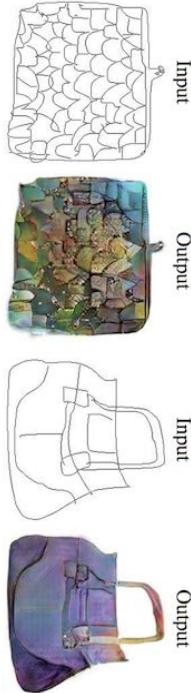
Brundage et al. "The Malicious Use of Artificial Intelligence: Forecasting, Prevention, and Mitigation" (2018)

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. **Image generation**
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

4. Image generation

Image-to-image translation



Isola et al. "Image-to-Image Translation with Conditional Adversarial Networks" (2016)

Text-to-image translation

this small bird has a pink breast and crown, and black primaries and secondaries.



Reed et al. "Generative Adversarial Text to Image Synthesis" (2016)



Introduction to Deep Learning - ICME Summer Workshops 2021

Source: Instagram, FaceApp - AI Face Editor

Deep learning examples

5. Text generation

GPT-3-aided poetry!

Human prompt (bolded):	GPT-3 output (with “several hours” of trial and error):
1. Optical character recognition (OCR)	
2. Image classification	
3. Text translation	
4. Image generation	
5. Text generation	
6. Audio generation	
7. Reinforcement learning	
8. Energy	
9. Healthcare	
10. Biology	
11. Development	
12. Investing	
13. Surveillance	

Source: <https://arr.am/2020/07/14/elon-musk-by-dr-seuss-gpt-3/> (2020)

Introduction to Deep Learning - ICME Summer Workshops 2021

...

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. Image generation
5. **Text generation**
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

5. Text generation

GPT-3-generated text:

Human prompt is bolded.

Movie plot summaries (from Wikipedia):

The Matrix (1999): "A computer hacker learns from mysterious rebels about the true nature of his reality and his role in the war against its controllers." (For the purposes of this post, we can set aside some of the specific details and just say that, basically, Neo learns that our world isn't real—that what we perceive isn't the real thing, but a simulation or “matrix” created to fool us.)

Inception (2010): "While competing with a rival thief to recover a potent artefact, Dom Cobb (Leonardo DiCaprio) falls into a dreamworld where he must ‘extract’ an idea from an influential businessman while being chased by the FBI."

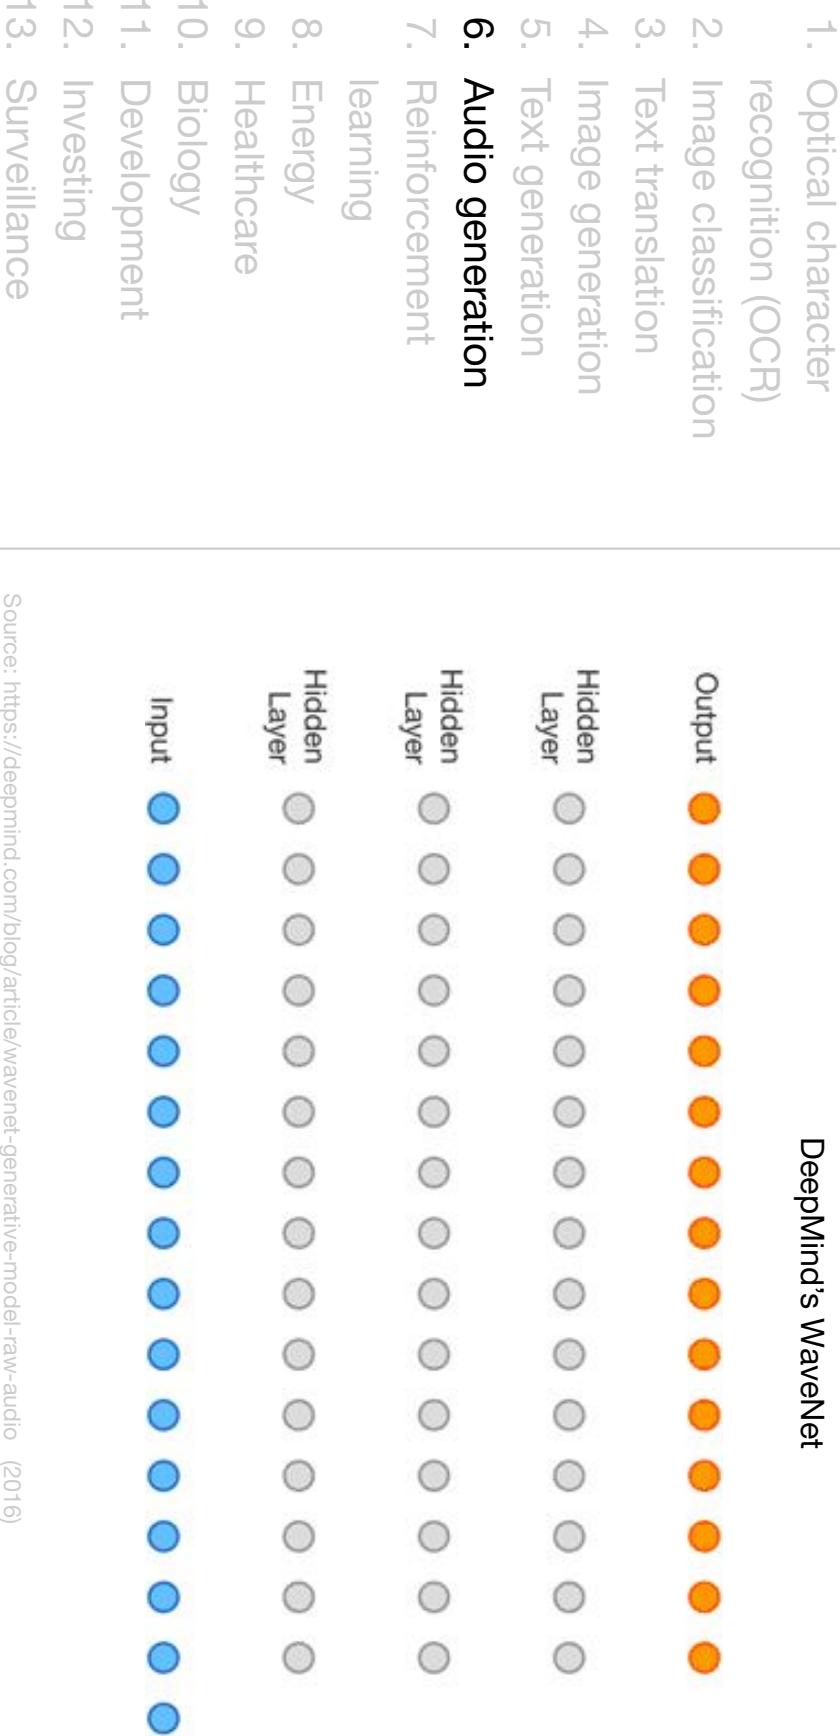
Source: <https://www.gwern.net/GPT-3> (2020)

Introduction to Deep Learning - ICME Summer Workshops 2021

Deep learning examples

6. Audio generation

DeepMind's WaveNet



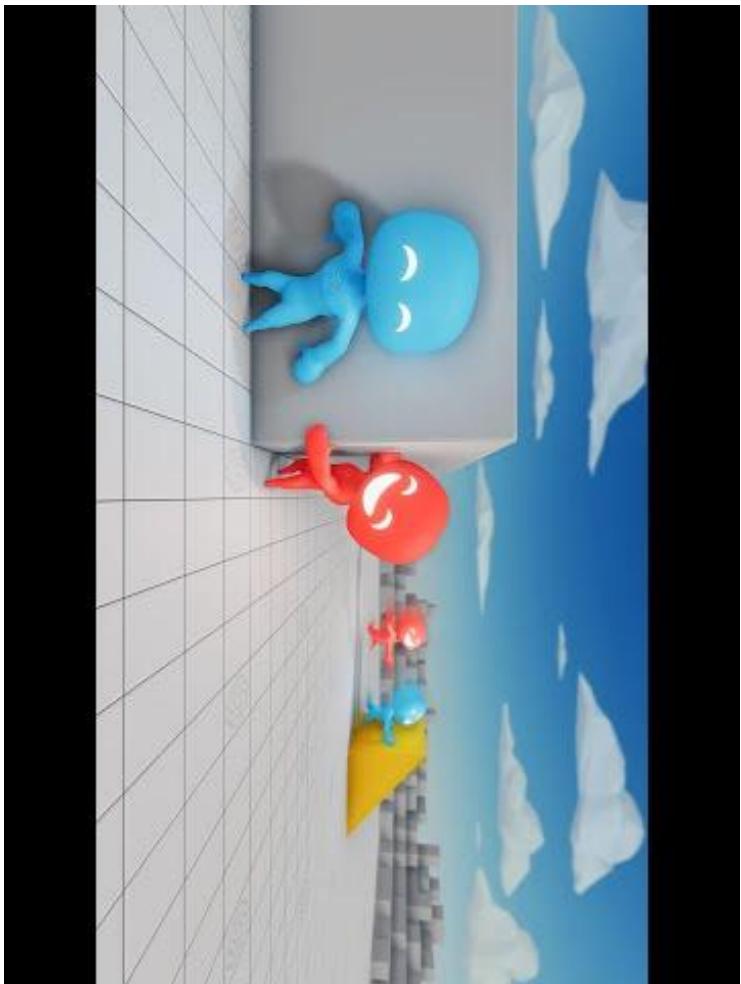
Source: <https://deepmind.com/blog/article/wavenet-generative-model-raw-audio> (2016)
Introduction to Deep Learning - ICME Summer Workshops 2021

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

7. Reinforcement learning

OpenAI's "Hide and Seek" Game



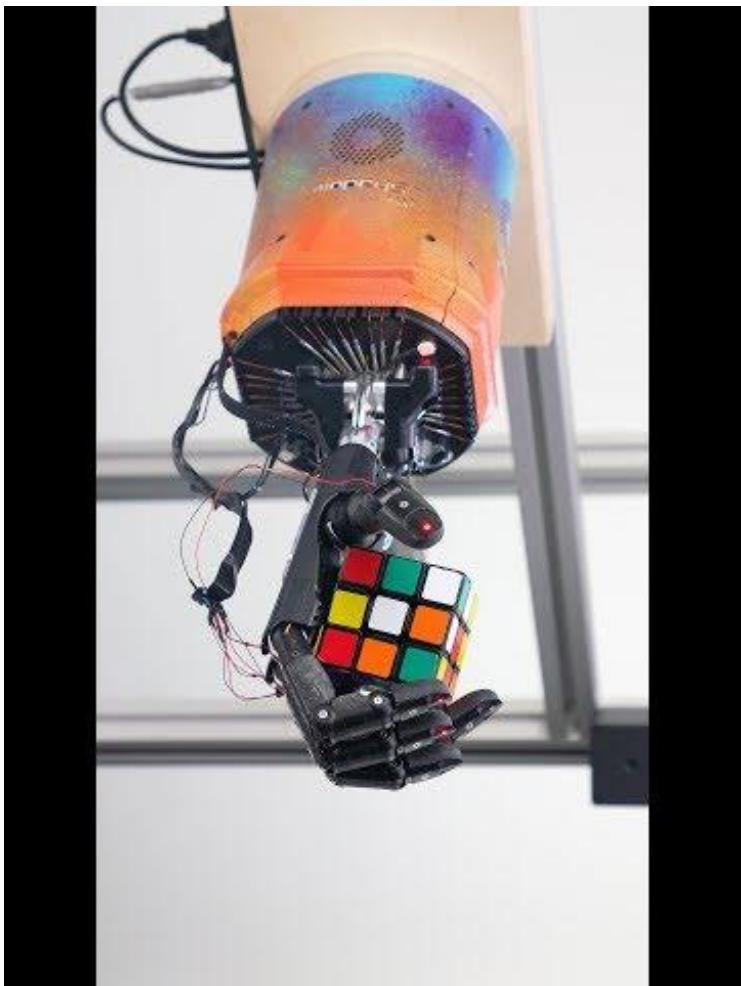
Source: <https://openai.com/blog/emergent-tool-use/> (2019)
Introduction to Deep Learning - ICME Summer Workshops 2021

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

7. Reinforcement learning

Solving a Rubik's cube with a robot hand



Source: <https://openai.com/blog/solving-rubiks-cube/> (2019)
Introduction to Deep Learning - ICME Summer Workshops 2021

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

7. Reinforcement learning

OpenAI vs. best human players in DOTA 2



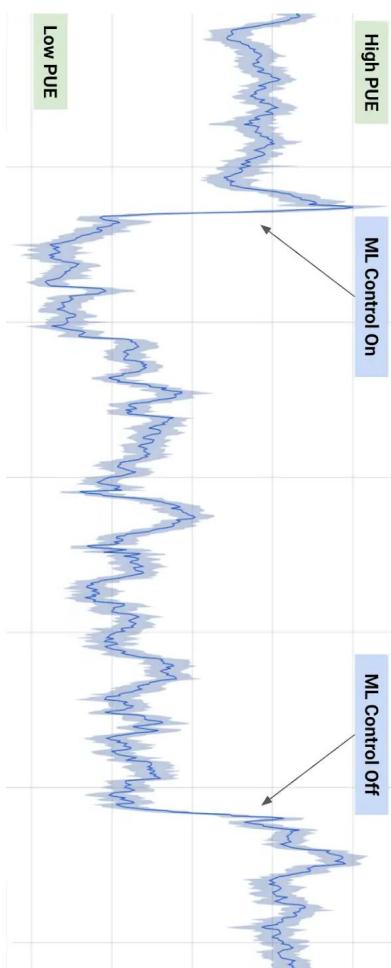
Source: <https://openai.com/projects/five/> (2019)
Introduction to Deep Learning - ICME Summer Workshops 2021

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

8. Energy

“DeepMind AI Reduces Google Data Centre Cooling Bill by 40%”



nature

Article | Published: 19 February 2020

Closed-loop optimization of fast-charging protocols for batteries with machine learning

Peter M. Atchia, Aditya Grover, Norman Jin, Kristen A. Severson, Todor M. Markov, Yang-Hung Liao, Michael H. Chen, Bryan Cheong, Nicholas Perkins, Zi Yang, Patrick K. Herring, Muratahan Aykol, Stephen J. Harris, Richard D. Braatz, Stefano Ermon & William C. Chueh

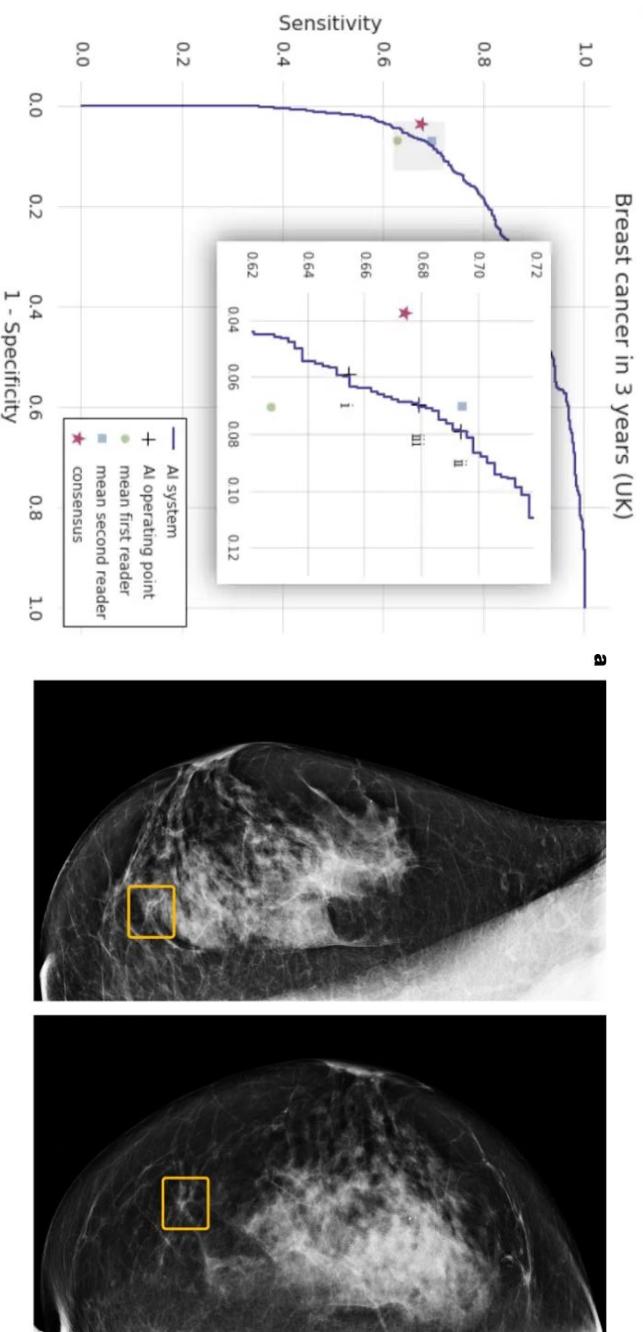
Source: (top) <https://deepmind.com/blog/article/deepmind-ai-reduces-google-data-centre-cooling-bill-40> (2016)
Introduction to Deep Learning - ICME Summer Workshops 2021

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy learning
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

9. Healthcare

DeepMind's AI for breast cancer screening



Source: McKinney et al. "International evaluation of an AI system for breast cancer screening" (2020)
Introduction to Deep Learning - ICME Summer Workshops 2021

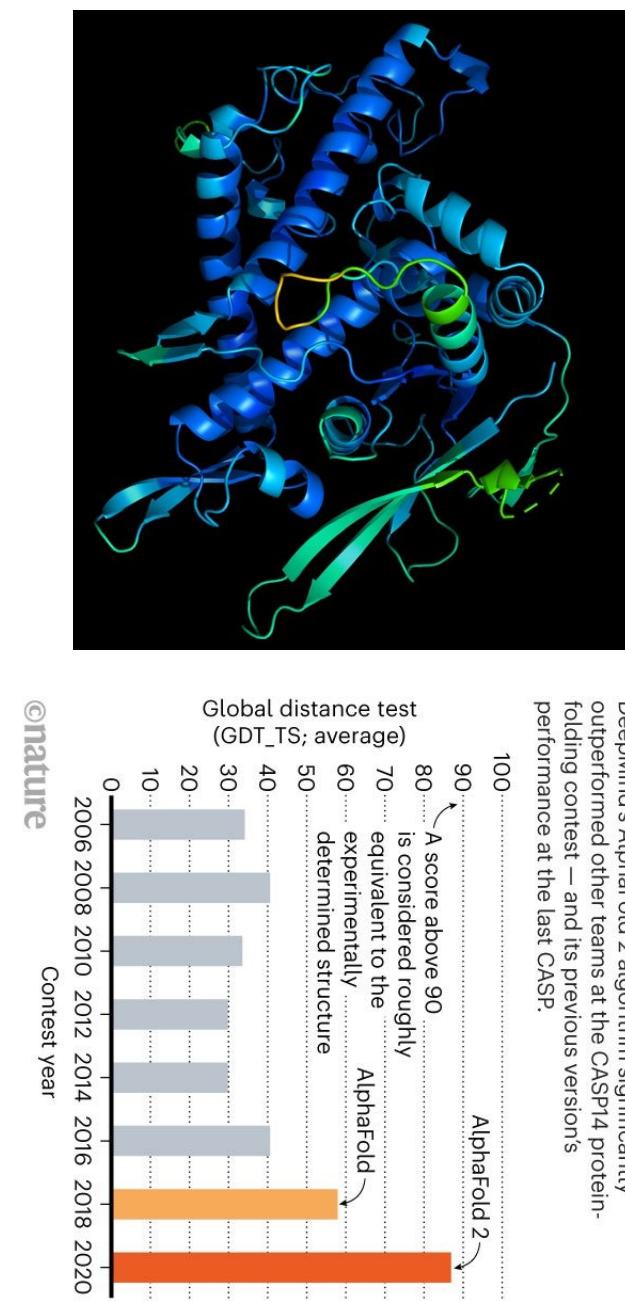
Deep learning examples

10. Biology

DeepMind's AlphaFold 2

STRUCTURE SOLVER

DeepMind's AlphaFold 2 algorithm significantly outperformed other teams at the CASP14 protein-folding contest – and its previous version's performance at the last CASP.



Source: "It will change everything": DeepMind's AI makes gigantic leap in solving protein structures" (2020)

Introduction to Deep Learning - ICME Summer Workshops 2021

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
- 11. Development**
12. Investing
13. Surveillance

11. Sustainable development

Estimating poverty using transfer learning and night lights



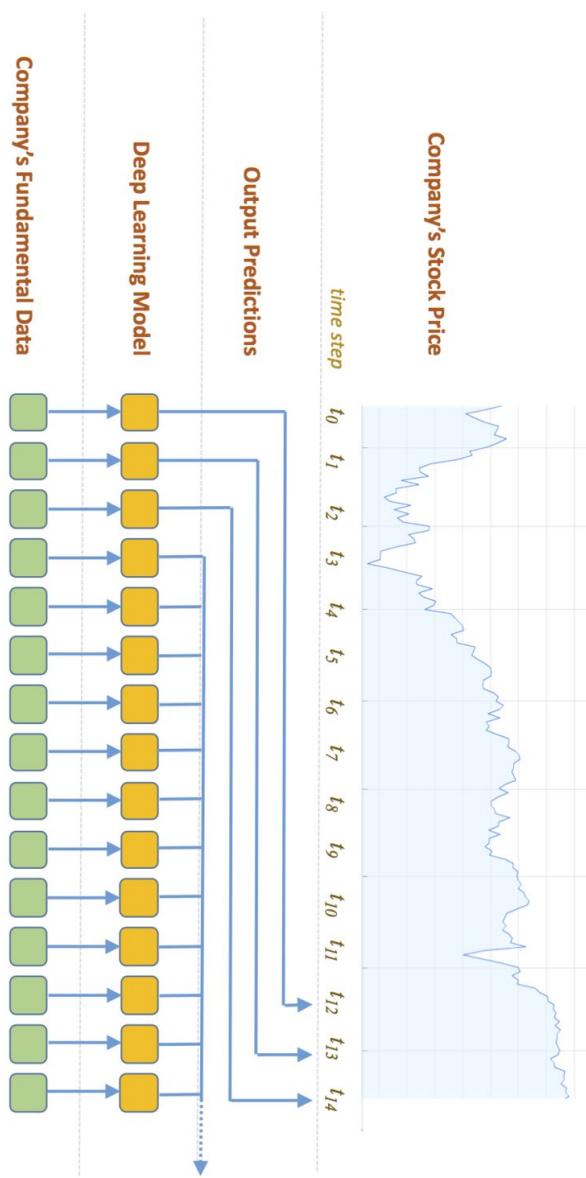
Source: Jean et al. "Combining satellite imagery and machine learning to predict poverty" (2016)
Introduction to Deep Learning - ICME Summer Workshops 2021

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
- 12. Investing**
13. Surveillance

12. Investing

Predicting stock price from company fundamentals



Source: Alberg and Lipton. "Improving Factor-Based Quantitative Investing by Forecasting Company Fundamentals" (2017)
Introduction to Deep Learning - ICME Summer Workshops 2021

Deep learning examples

1. Optical character recognition (OCR)
2. Image classification
3. Text translation
4. Image generation
5. Text generation
6. Audio generation
7. Reinforcement learning
8. Energy
9. Healthcare
10. Biology
11. Development
12. Investing
13. Surveillance

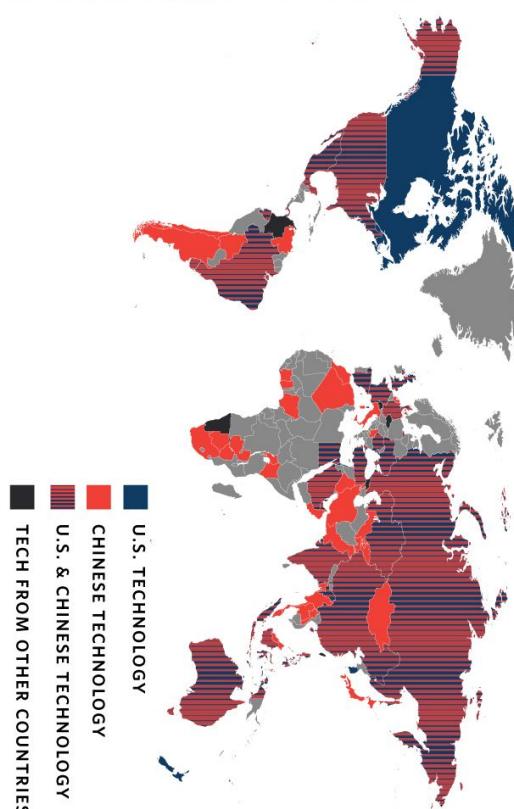


Attendees interacting with a facial recognition demonstration at this year's CES in Las Vegas. Joe Buglewicz for The New York Times

13. Surveillance

The New York Times
San Francisco Bans Facial Recognition Technology

MAP 1
AI Surveillance Technology Origin



Sources: <https://www.nytimes.com/2019/05/14/us/facial-recognition-ban-san-francisco.html> (2019)
<https://carnegieendowment.org/2019/09/17/global-expansion-of-ai-surveillance-pub-79847> (2019)

Introduction to Deep Learning - ICME Summer Workshops 2021

Why is deep learning at the forefront now?

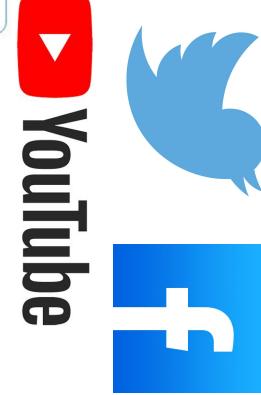
Big data



Compute



Microsoft
Azure



Introduction to Deep Learning - ICME Summer Workshops 2021

The collage includes:

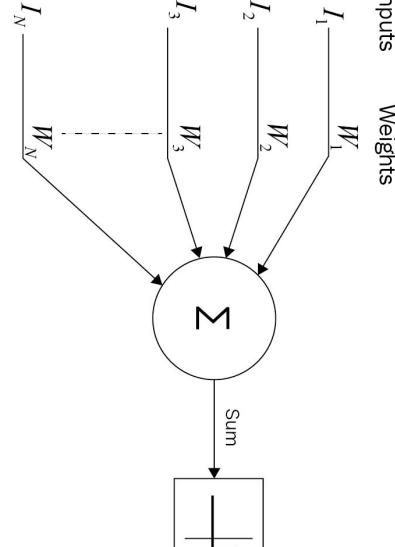
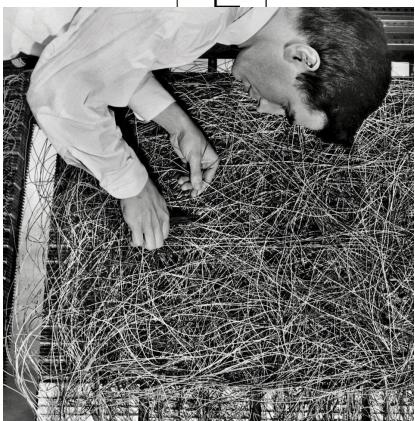
- An Apple Watch displaying a circular interface.
- A smartphone screen showing a home screen with various app icons.
- A screenshot of the Amazon.com website under the heading "Recommended for You". It shows book recommendations based on previous purchases, with titles like "Google Apps", "Google Tips", "Google APIs", and "Google.pedia".
- A screenshot of a financial chart for Tesla (TSLA) stock, showing a price drop from 252.39 to 251.74.
- A screenshot of a server rack with many blue LED lights.
- A photograph of a Tesla GPU card.

Why is deep learning at the forefront now?

Algorithms

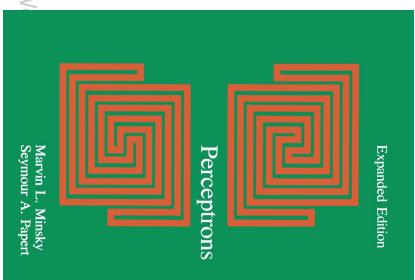


1943: McCulloch and Pitts develop a mathematical model of the neuron



1957: Rosenblatt invents the perceptron algorithm and machine

1969: Minsky and Papert improve the perceptron and prove it cannot model nonlinearities like XOR



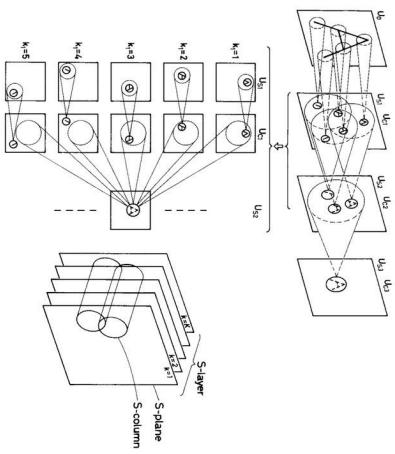
1974–1980:
The First AI Winter



Why is deep learning at the forefront now?

Algorithms

1980: Fukushima creates the “neocognitron”, introducing convolutional and downsampling layers

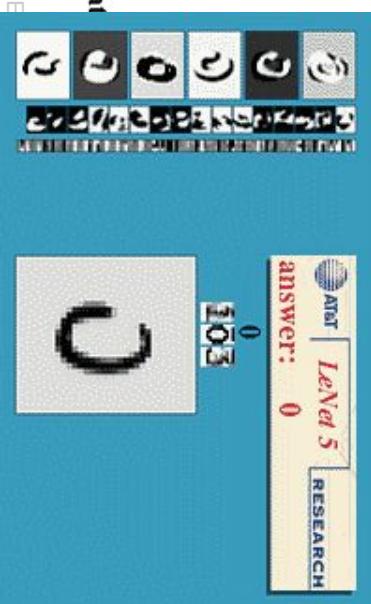


1986: Rumelhart, Hinton, and Williams apply backpropagation to train neural networks



D.E. Rumelhart, G.E. Hinton, R.J. Williams
Learning representation by back-propagation errors. *Nature*, 323 (1986), pp. 533–536

1989: LeCun uses backprop to train convolutional neural networks to read digits



1987–1994:
The Second AI Winter



Why is deep learning at the forefront now?

Algorithms



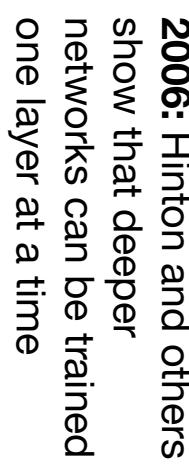
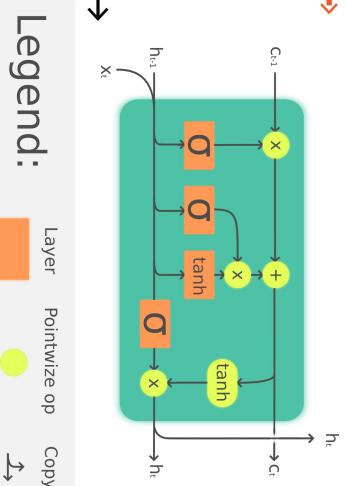
1990s: Computers become faster, GPUs developed



1991: The vanishing gradient problem is identified



1997: Hochreiter and Schmidhuber develop long short-term memory



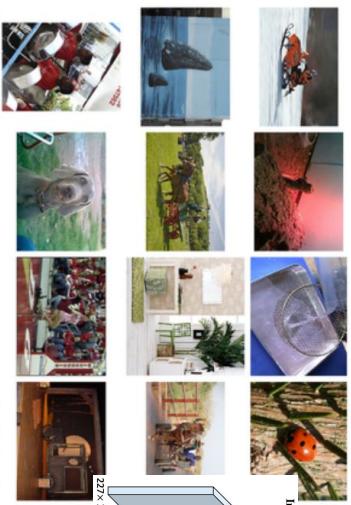
2006: Hinton and others show that deeper networks can be trained one layer at a time

Why is deep learning at the forefront now?

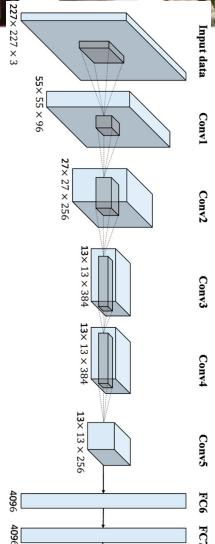
Algorithms



2009: Fei-fei Li releases ImageNet, a classification dataset of 14 million labeled images



2012: AlexNet, a CNN, wins the Large Scale Visual Recognition Challenge – by a lot



2014: Goodfellow designs generative adversarial networks for data generation



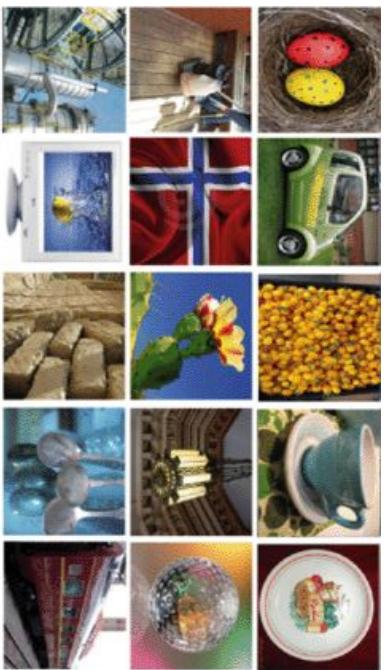
2015-2016: Deep learning frameworks like TensorFlow, Keras, and PyTorch are released



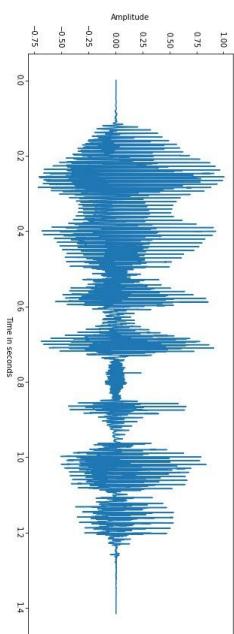
What's so special about deep learning?

Deep learning has made significant improvements in performance on tasks involving unstructured data

Images



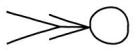
Speech

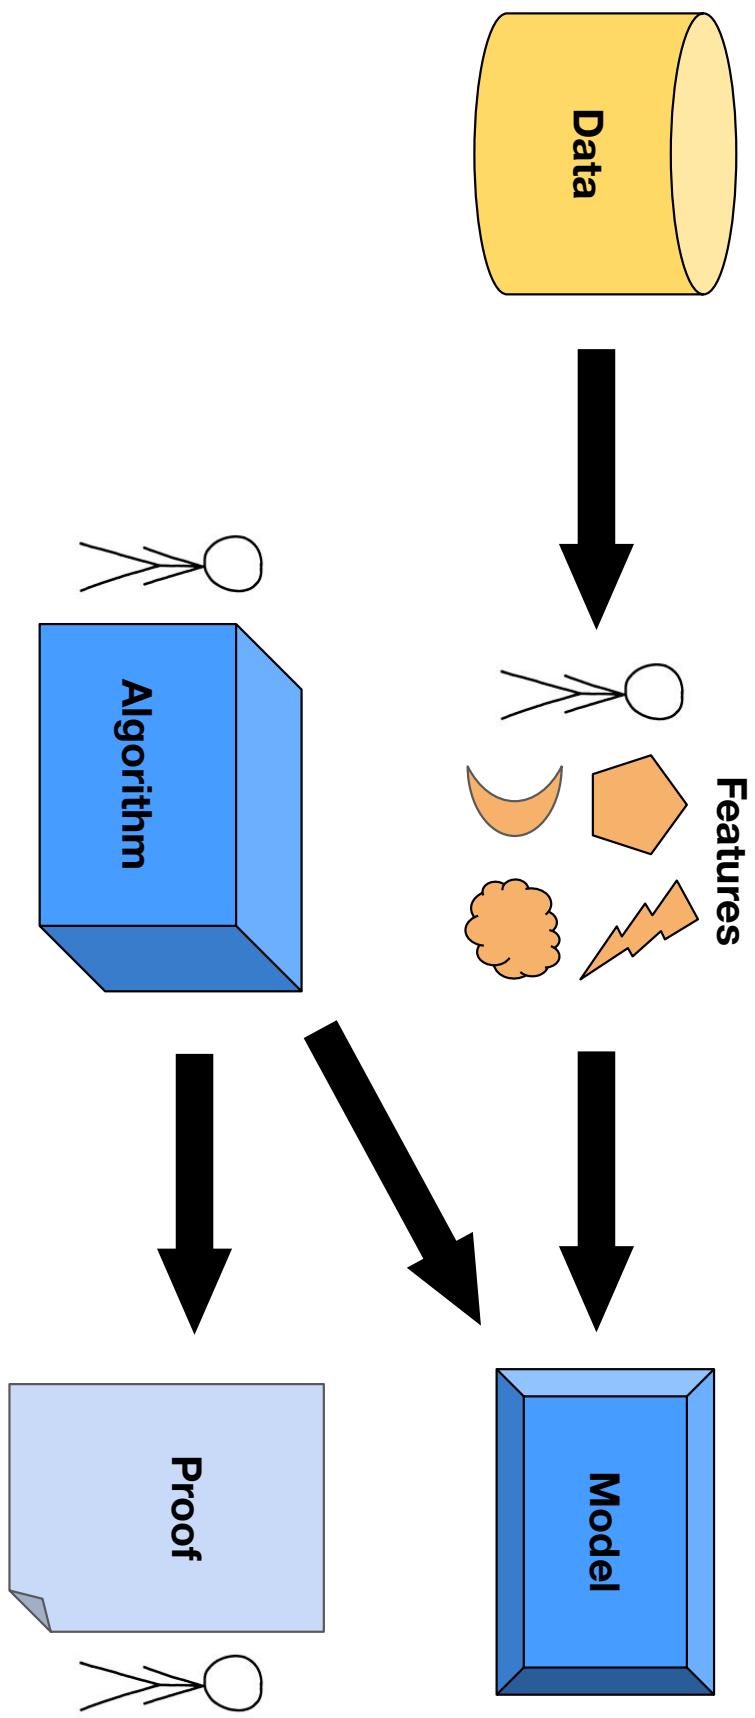


Text

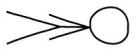
It is a truth universally acknowledged, that a single man in possession of a good fortune, must be in want of a wife. However little known the feelings or views of such a man may be on his first entering a neighbourhood, this truth is so well fixed in the minds of the surrounding families, that he is considered ...

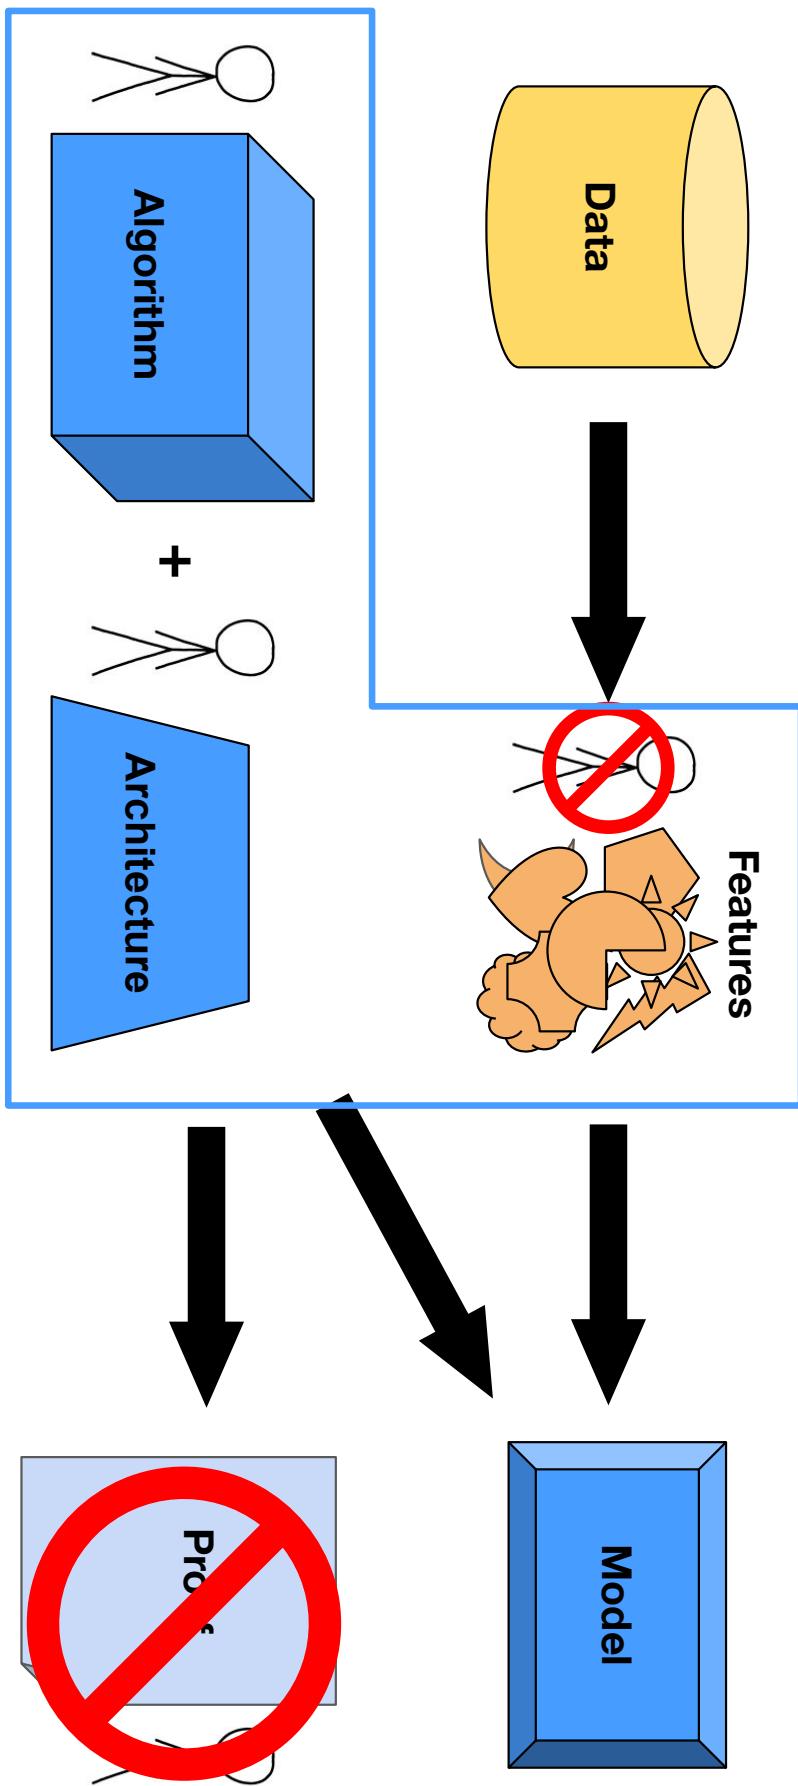
Traditional machine learning

 = Human involved
in designing



Deep learning

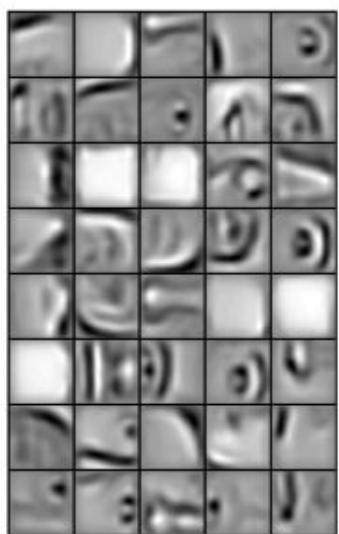
 = Human involved
in designing



Introduction to Deep Learning - ICME Summer Workshops 2021

Deep learning-generated features

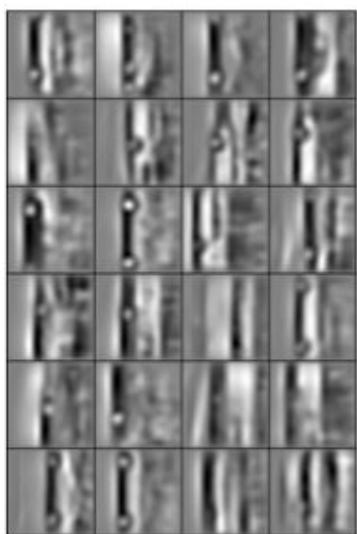
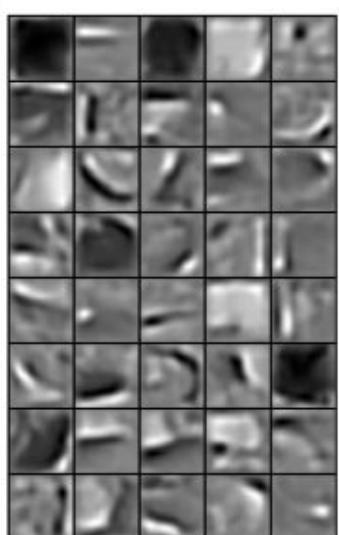
faces



cars



elephants

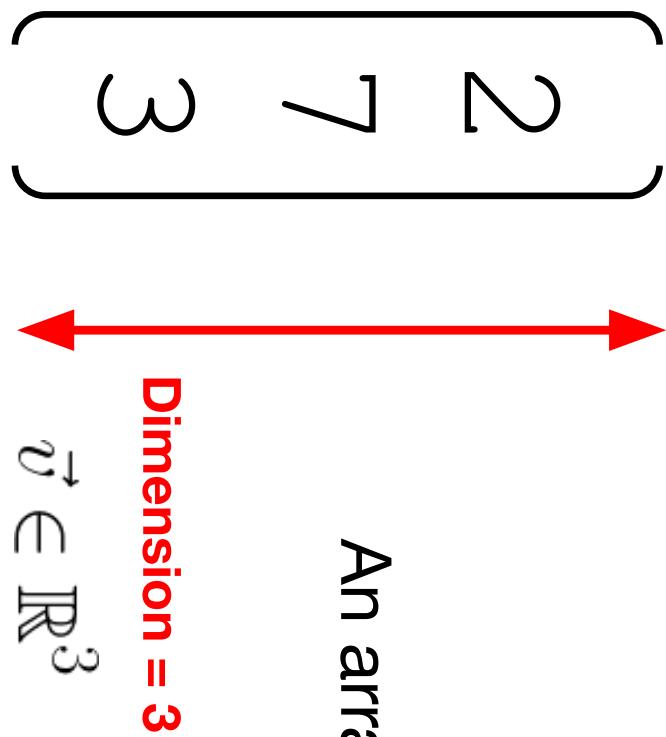


Math Review

Vectors, matrices, and tensors

What is a vector?

An array (or list) of numbers

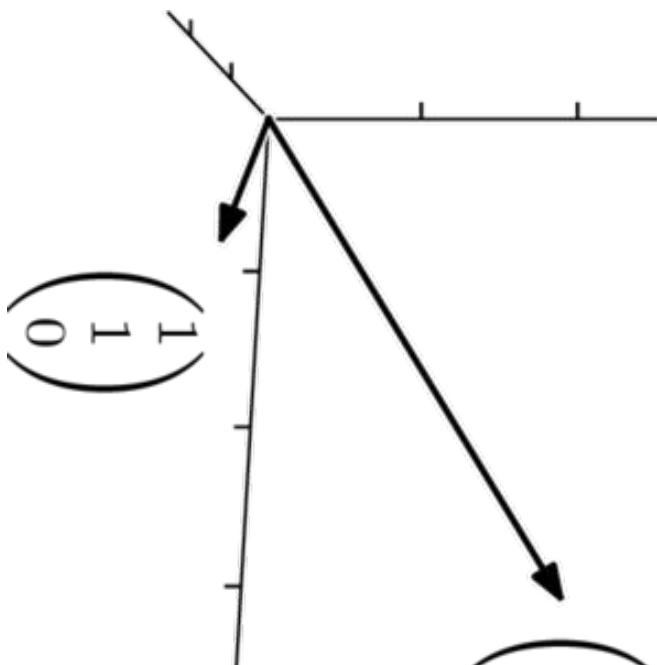


Dimension = 3

$$\vec{v} \in \mathbb{R}^3$$

What is a vector?

$$\begin{pmatrix} 0 \\ 3 \\ 2 \end{pmatrix}$$



In addition to linear algebra, you may have learned about vectors in physics courses as entities with direction and magnitude

Here are two more vectors of dimension 3, visualized in Euclidean space

Vector operations

1. Scalar multiplication

$$4 * \begin{bmatrix} 2 \\ 7 \\ 3 \end{bmatrix} = \begin{bmatrix} 8 \\ 28 \\ 12 \end{bmatrix}$$

Vector operations

2. Addition

$$\begin{bmatrix} 3 \\ 7 \end{bmatrix} + \begin{bmatrix} 6 \\ 3 \end{bmatrix} = \begin{bmatrix} 9 \\ 10 \end{bmatrix}$$

Vector operations

3. Transpose

$$\begin{matrix} 3 & 7 & 2 \end{matrix}^T = \begin{matrix} 2 \\ 7 \\ 3 \end{matrix}$$

A diagram illustrating the transpose operation. A horizontal row vector $\begin{matrix} 3 & 7 & 2 \end{matrix}$ is enclosed in a blue oval. A blue arrow points from this oval to the right, followed by an equals sign. To the right of the equals sign is a red arrow pointing down, indicating the transformation to a vertical column vector. The column vector $\begin{matrix} 2 \\ 7 \\ 3 \end{matrix}$ is enclosed in a red oval.

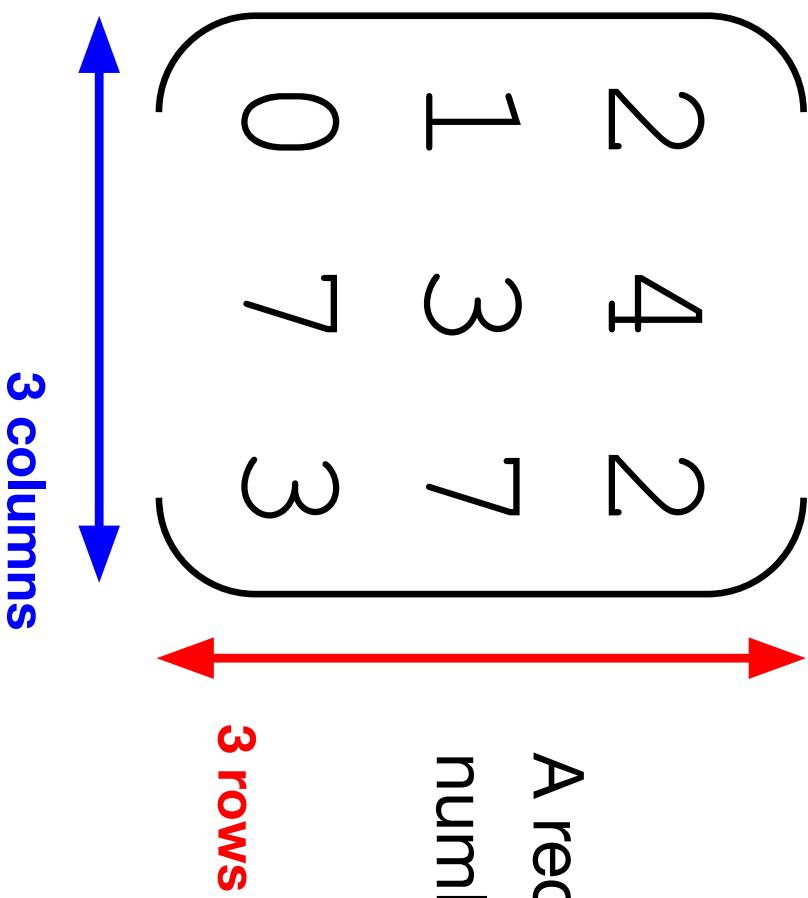
Vector operations

4. Inner product

$$\begin{bmatrix} 2 & 7 & 3 \end{bmatrix} \begin{bmatrix} 1 \\ 3 \\ 6 \end{bmatrix} = 2 * 1 + 7 * 3 + 3 * 6 = 41$$

$$\vec{u}^\top \vec{v} = w$$

What is a matrix?



A rectangular array (or table) of numbers

What is a matrix?

$$\mathbf{A} \in \mathbb{R}^{m \times n}$$
$$\begin{bmatrix} 1 & 2 & \cdots & n \\ a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ a_{31} & a_{32} & \cdots & a_{3n} \\ \vdots & \vdots & \ddots & \vdots \\ \vdots & \vdots & \ddots & \vdots \\ m & a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}$$

What is a matrix?

$$\begin{pmatrix} 2 \\ 1 \\ 0 \end{pmatrix} \begin{pmatrix} 4 \\ 3 \\ 7 \end{pmatrix} \begin{pmatrix} 2 \\ 7 \\ 3 \end{pmatrix}$$

You can think of a matrix as a row vector of column vectors

What is a matrix?

$$\begin{pmatrix} 2 & 4 & 2 \\ 1 & 3 & 7 \\ 0 & 7 & 3 \end{pmatrix}$$

... or as a column vector of row
vectors

Matrix operations

1. Scalar multiplication
2. Addition

$$2 * \begin{pmatrix} 2 & 4 & 2 \\ 1 & 3 & 7 \\ 0 & 7 & 3 \end{pmatrix} = \begin{pmatrix} 4 & 8 & 4 \\ 2 & 6 & 14 \\ 0 & 14 & 6 \end{pmatrix}$$

Matrix operations

3. Transpose

A

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{bmatrix}$$

Source: <https://en.wikipedia.org/wiki/Transpose>

Introduction to Deep Learning - ICME Summer Workshops 2021

Matrix operations

3. Transpose

$$\begin{pmatrix} 2 \\ 1 \\ 0 \\ 7 \\ 3 \end{pmatrix}^T = \begin{pmatrix} 2 & 1 & 0 \\ 4 & 3 & 7 \\ 7 & 3 & 2 \end{pmatrix}$$

Matrix operations

4. Multiplication with a vector

$$\begin{pmatrix} 2 & 4 & 2 \\ 1 & 3 & 7 \\ 0 & 7 & 3 \end{pmatrix} \begin{pmatrix} 1 \\ 3 \\ 2 \end{pmatrix} = \begin{pmatrix} 18 \\ 24 \\ 27 \end{pmatrix}$$

Matrix operations

4. Multiplication with a vector

$$\begin{pmatrix} 2 \\ 1 \\ 0 \\ 7 \\ 3 \end{pmatrix} \begin{pmatrix} 4 \\ 3 \\ 7 \\ 1 \\ 2 \end{pmatrix} = \begin{pmatrix} 18 \\ 24 \\ 27 \end{pmatrix}$$

Matrix operations

4. Multiplication with a vector

$$1 \begin{pmatrix} 2 \\ 1 \\ 0 \end{pmatrix} + 3 \begin{pmatrix} 4 \\ 3 \\ 7 \end{pmatrix} + 2 \begin{pmatrix} 2 \\ 7 \\ 3 \end{pmatrix} = \begin{pmatrix} 18 \\ 24 \\ 27 \end{pmatrix}$$

Matrix operations

4. Multiplication with a vector

$$\begin{pmatrix} 2 & 4 & 2 \\ 1 & 3 & 7 \\ 0 & 7 & 3 \end{pmatrix} \begin{pmatrix} 1 \\ 3 \\ 2 \end{pmatrix} = \begin{pmatrix} 18 \\ 24 \\ 27 \end{pmatrix}$$

Matrix operations

4. Multiplication with a vector

$$\begin{bmatrix} 2 & 4 & 2 \end{bmatrix} \begin{bmatrix} 1 \\ 3 \\ 2 \end{bmatrix} = \begin{bmatrix} 18 \\ 24 \\ 27 \end{bmatrix}$$

Matrix operations

4. Multiplication with a vector

$$\vec{b} = \begin{pmatrix} \vec{a}_1^\top \\ \vec{a}_2^\top \\ \vec{a}_3^\top \end{pmatrix}$$

Matrix operations

4. Multiplication with a vector

$$\begin{pmatrix} \vec{a}_1^\top \vec{b} \\ \vec{a}_2^\top \vec{b} \\ \vec{a}_3^\top \vec{b} \end{pmatrix} = \begin{pmatrix} 1 \\ 8 \\ 2 \\ 4 \\ 2 \\ 7 \end{pmatrix}$$

Matrix operations

4. Multiplication with another matrix

$$\begin{pmatrix} 2 & 4 & 2 \\ 1 & 3 & 7 \\ 0 & 7 & 3 \end{pmatrix} \begin{pmatrix} 1 & 9 \\ 2 & 3 \\ 4 & 0 \end{pmatrix} = \begin{pmatrix} 18 & 26 \\ 24 & 37 \\ 27 & 12 \end{pmatrix}$$

Diagram illustrating matrix multiplication:

- The first matrix has 3 columns (labeled "3 columns" with blue double-headed arrows) and 3 rows (labeled "3 rows" with red double-headed arrows).
- The second matrix has 2 columns (labeled "2 columns" with blue double-headed arrows) and 3 rows (labeled "3 rows" with red double-headed arrows).
- The result matrix has 2 columns (labeled "2 columns" with blue double-headed arrows) and 2 rows (labeled "3 rows" with red double-headed arrows).

Matrix operations

4. Multiplication with another matrix

$$\begin{pmatrix} 2 & 4 & 2 \\ 1 & 3 & 7 \\ 0 & 7 & 3 \end{pmatrix} \begin{pmatrix} 1 \\ 3 \\ 0 \end{pmatrix} = \begin{pmatrix} 18 & 26 \\ 24 & 37 \\ 27 & 12 \end{pmatrix}$$

Matrix operations

4. Multiplication with another matrix

$$A \begin{pmatrix} 1 \\ 3 \\ 2 \\ 4 \end{pmatrix} = \begin{pmatrix} A \\ 3 \\ 2 \end{pmatrix} A \begin{pmatrix} 9 \\ 0 \\ 4 \end{pmatrix}$$

Matrix operations

4. Multiplication with another matrix

$$\begin{array}{ccc} \vec{a}_1 & \rightarrow & \begin{bmatrix} 1 & 7 \\ 2 & 4 \end{bmatrix} \cdot \begin{bmatrix} 3 & 3 \\ 5 & 2 \end{bmatrix} = \begin{bmatrix} \vec{a}_1 \cdot \vec{b}_1 & \vec{a}_1 \cdot \vec{b}_2 \\ \vec{a}_2 \cdot \vec{b}_1 & \vec{a}_2 \cdot \vec{b}_2 \end{bmatrix} \\ \vec{a}_2 & \rightarrow & \downarrow \quad \downarrow \\ A & B & C \end{array}$$

Source: https://ml-cheatsheet.readthedocs.io/en/latest/linear_algebra.html#matrix-multiplication

Introduction to Deep Learning - ICME Summer Workshops 2021

Poll:

Matrix multiplication

Go to tinyurl.com/dlworkshop2021

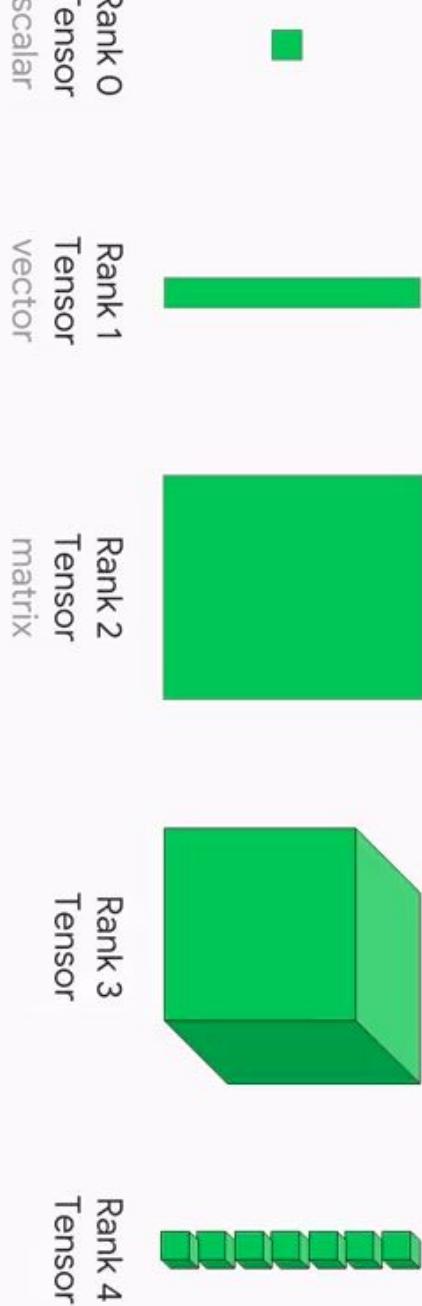
What are the dimensions of the matrix resulting from:

$$\begin{pmatrix} 2 & 4 & 2 & 8 & 4 & 1 & 5 & 0 \\ 1 & 3 & 7 & 9 & 3 & 2 & 6 & 8 \\ 0 & 7 & 3 & 1 & 0 & 4 & 9 & 3 \end{pmatrix} \cdot \begin{pmatrix} 3 & 4 \\ 0 & 2 \\ 2 & 9 \\ 5 & 5 \\ 3 & 7 \\ 6 & 4 \\ 2 & 1 \\ 4 & 0 \end{pmatrix}$$

- A. 8×8
- B. 8×2
- C. 3×2
- D. Not valid

What is a tensor?

A tensor is an N-dimensional array of data



Source: <https://medium.com/mlait/tensors-representation-of-data-in-neural-networks-bbe8a711b93b>

Introduction to Deep Learning - ICME Summer Workshops 2021

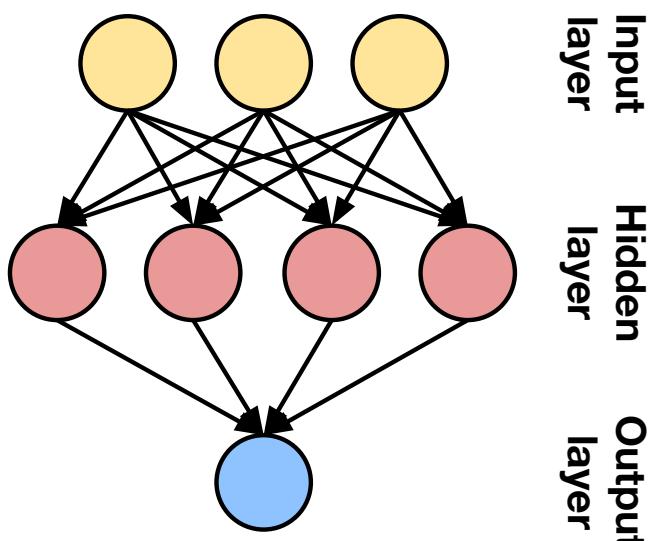
Neural Networks

Neural networks

We'll start with a *fully connected* neural network.

1. Single layer
2. Single neuron
3. Multiple layers
4. Input and output

Later, we'll cover other types of layers (convolution, max-pool, ...)



What does a single layer in a neural network do?

First, it involves some matrix operations...

$$\begin{pmatrix} 2 & 4 & 2 \\ 1 & 3 & 7 \\ 0 & 7 & 3 \end{pmatrix} \begin{pmatrix} 1 \\ 3 \\ 0 \end{pmatrix} + \begin{pmatrix} -3 \\ 2 \\ 2 \end{pmatrix} = \begin{pmatrix} 15 \\ 24 \\ 29 \end{pmatrix}$$

What does a single layer in a neural network do?

Then, the intermediate values are passed through a nonlinear function...

The diagram illustrates the application of a nonlinear activation function g to a vector of inputs. On the left, a red arrow labeled "Activation function" points from a red oval containing the letter g towards a vector of inputs. This vector is enclosed in large black parentheses and contains three elements: 15, 24, and 29. An equals sign follows the vector. To the right of the equals sign is another vector, also enclosed in large black parentheses, containing three elements: $g(15)$, $g(24)$, and $g(29)$.

$$\begin{pmatrix} 15 \\ 24 \\ 29 \end{pmatrix} = \begin{pmatrix} g(15) \\ g(24) \\ g(29) \end{pmatrix}$$

Activation functions

There exist a number of activation functions, some of which were designed with neurons in mind

What's the purpose?

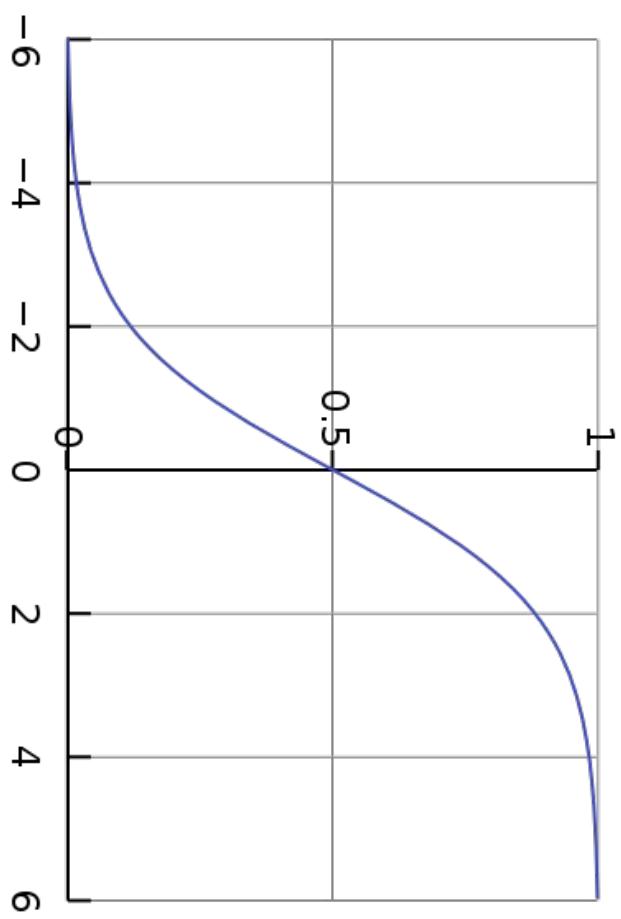
- Most phenomena are not linear, so we introduce some nonlinearities to the network

Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0. \end{cases}$	Perception variant	
Sign (Signum)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Linear	$\phi(z) = z$	Adaline, linear regression	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	
Hyperbolic tangent	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multi-layer Neural Networks	
Rectifier, ReLU (Rectified Linear Unit)	$\phi(z) = \max(0, z)$	Multi-layer Neural Networks	
Rectifier, softplus	$\phi(z) = \ln(1 + e^z)$	Multi-layer Neural Networks	

<sup>“A visual proof that neural nets can compute any function”:
<http://neuralnetworksanddeeplearning.com/chap4.html></sup>

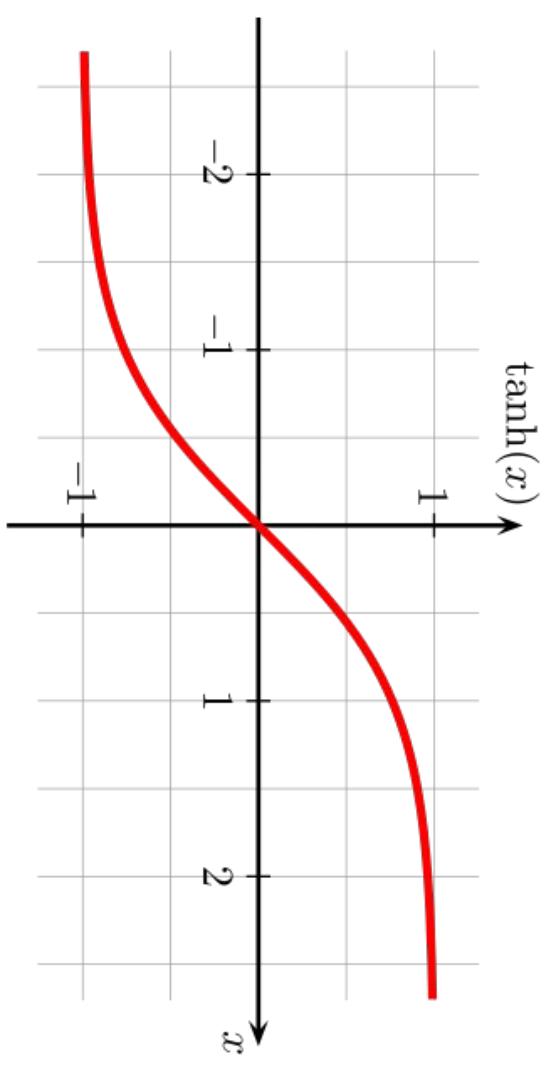
Sigmoid (logistic) activation function

$$\text{sigmoid}(x) = \frac{1}{1 + e^{-x}}$$



Hyperbolic tangent activation function

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

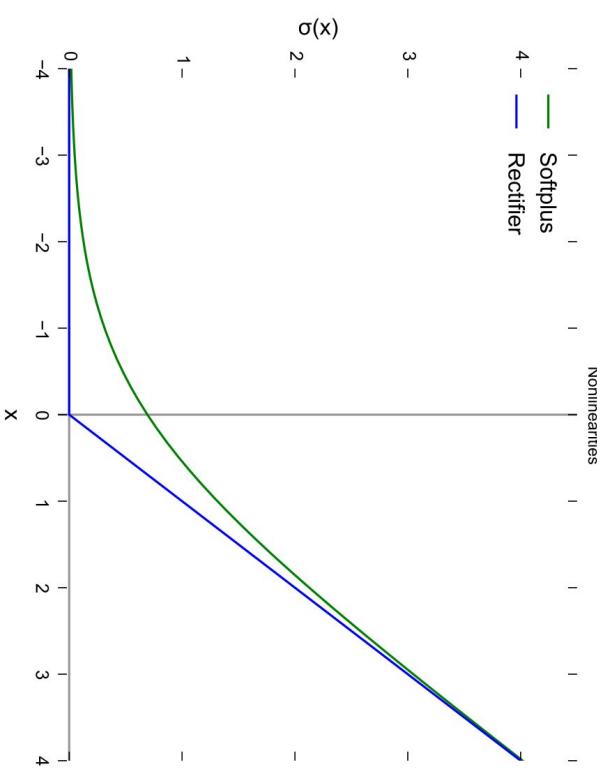


ReLU activation function

Stands for Rectified Linear Unit

$$\text{ReLU}(x) = \max(0, x)$$

A smooth version of ReLU is softplus



$$\text{softplus}(x) = \ln(1 + e^x)$$

Poll:

Activation functions

Go to tinyurl.com/dlworkshop2021

As its input goes to infinity, the output of the logistic function goes to...

$$\lim_{x \rightarrow \infty} \frac{1}{1 + e^{-x}} =$$

A. 0

B. 0.5

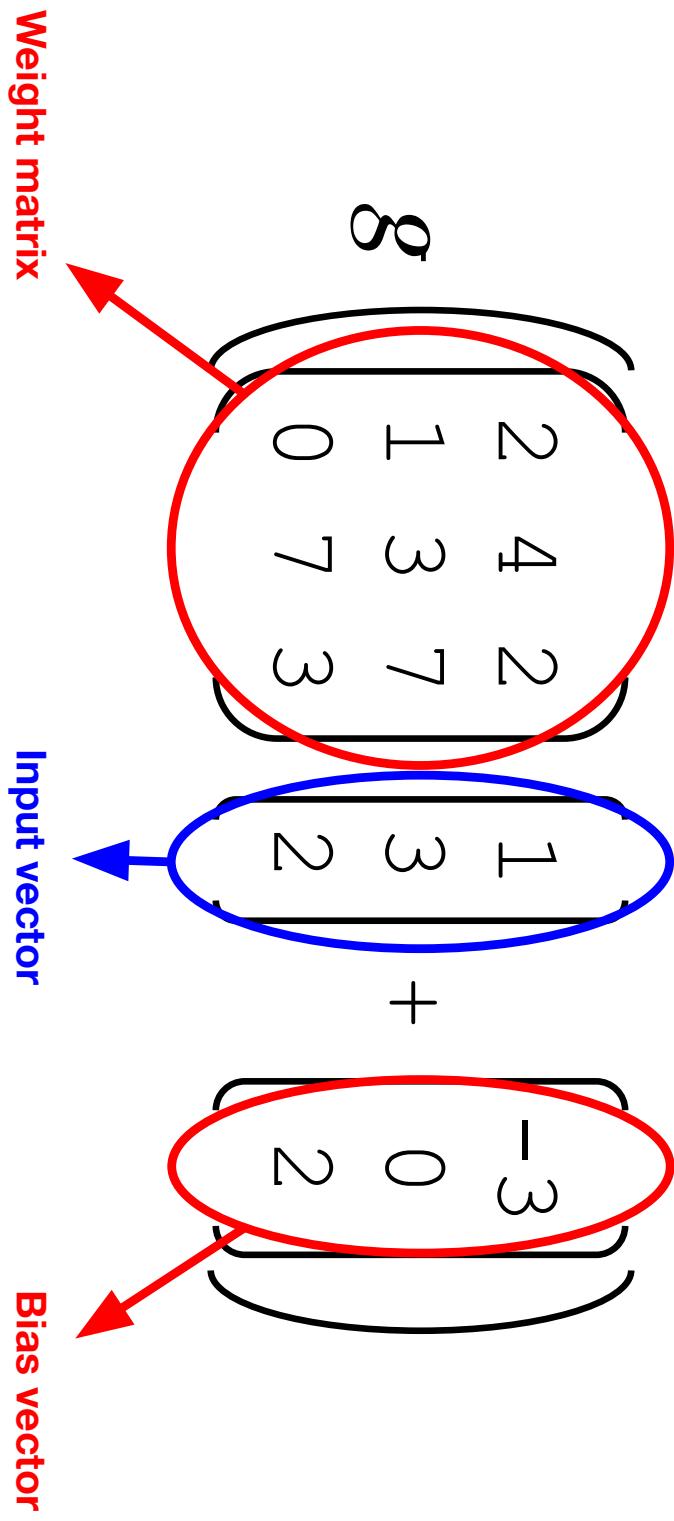
C. 1

D. Infinity

Activation functions FAQ

1. Why these activation functions? Why not x^2 or other nonlinear functions?
 - a. History of the field
 - b. Normalizing layer outputs, concerns with backpropagation (to be discussed)
2. How does one choose among the various options for activation functions?
 - a. You don't have to choose -- architectures have already been designed
 - b. Sigmoid, tanh, and ReLU are popular choices, with ReLU being the most common. ReLU has fewer vanishing gradient problems than the other 2 (see Session 3). If you're developing your own architecture, you should use your validation set to experiment with choice of activation function.

Recall: a single layer is matrix multiplication followed by a nonlinearity



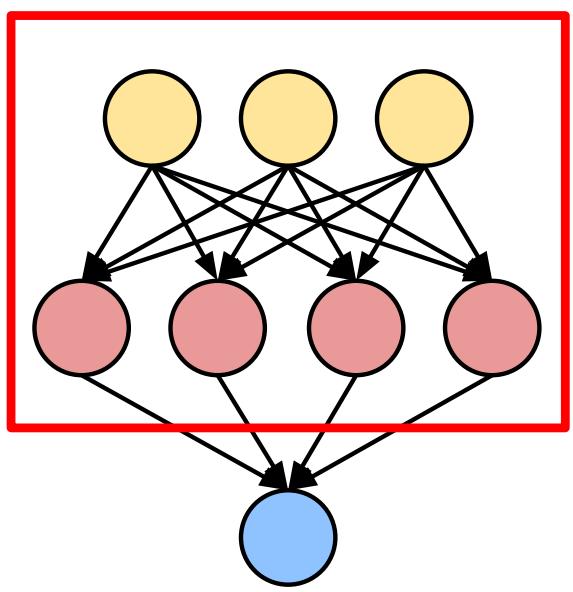
What are the matrix and vectors?

$$g = \begin{pmatrix} 2 & 4 & 2 \\ 1 & 3 & 7 \\ 0 & 7 & 3 \end{pmatrix} \begin{pmatrix} 1 \\ 3 \\ 2 \end{pmatrix} + \begin{pmatrix} -3 \\ 0 \\ 2 \end{pmatrix}$$

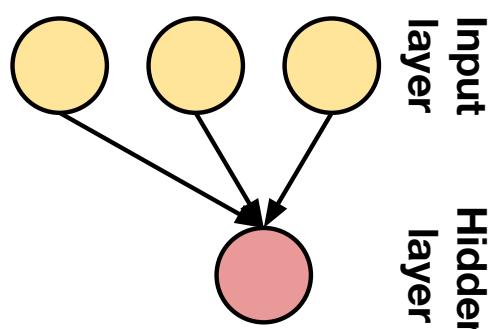
Trainable parameters

A more compact notation

$$g(\mathbf{W}\vec{x} + \vec{b})$$



What does a single neuron do?



Input layer
Hidden layer

Each circle is one element of an input or hidden layer

Each edge is a weight

$$g(\vec{w}_1^\top \vec{x} + b_1)$$

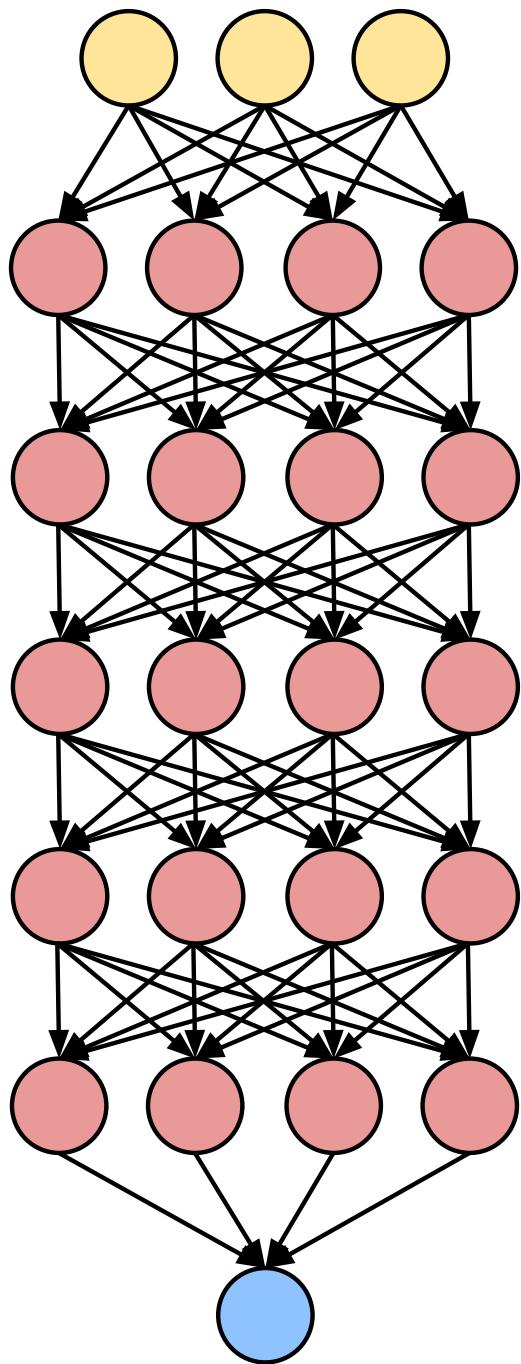
Weights for each neuron

$$\mathbf{W} = \begin{pmatrix} 2 & 4 & 2 \\ 1 & 3 & 7 \\ 0 & 7 & 3 \end{pmatrix} = \begin{pmatrix} w_1^\top \\ w_2^\top \\ w_3^\top \end{pmatrix}$$

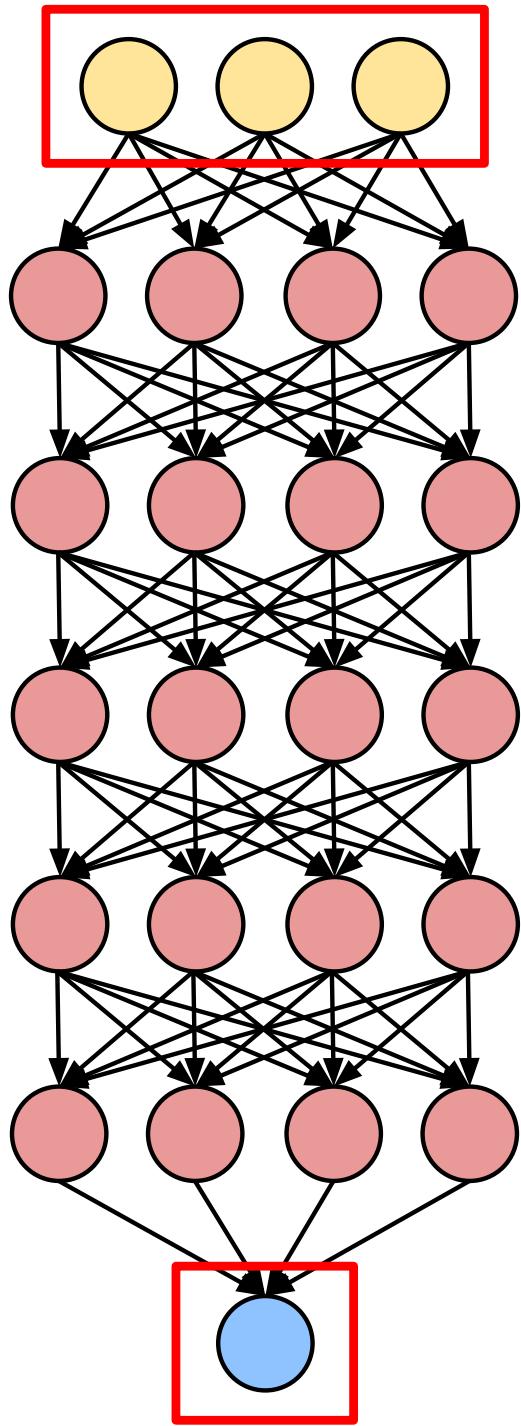
A fully connected layer with weight matrix $\mathbf{W} \in \mathbb{R}^{m \times n}$ is a function from \mathbb{R}^n to \mathbb{R}^m

Stacking multiple layers

It becomes a *deep* neural network when you use many layers



What are the inputs and outputs?



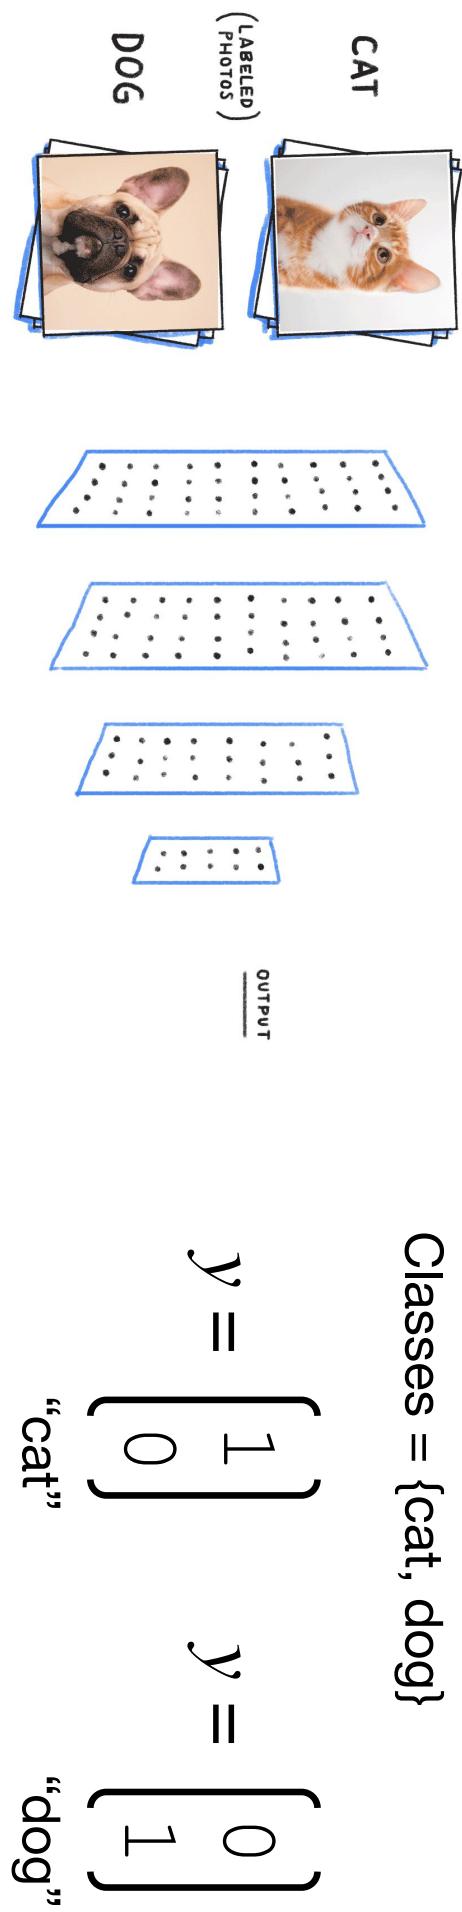
Defining the problem

- You want to predict some output \mathbf{y} given input \mathbf{x}
- If \mathbf{y} is a discrete output (e.g. image classes), then it's a *classification problem*
- If \mathbf{y} is a continuous variable (e.g. price), then it's a *regression problem*

$$y = \text{DNN}(\vec{x})$$

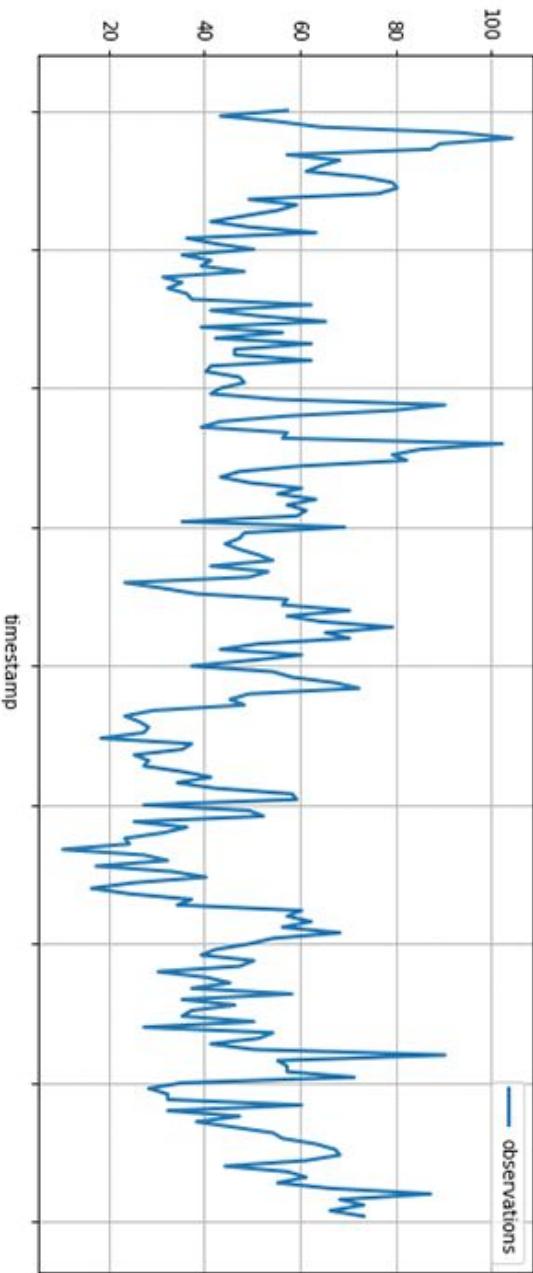
Defining the problem

- If y is a discrete output, you can use **one-hot encoding**
- If y is a continuous variable, it can be output by the network directly



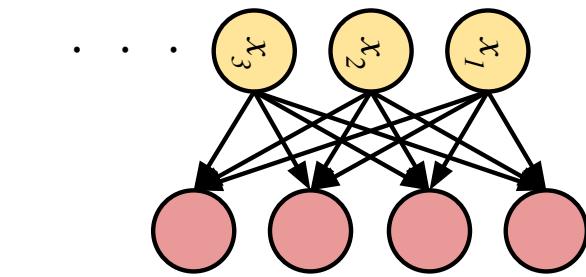
Defining the problem

- If x is a discrete output, you can use **one-hot encoding**
- If x is a continuous variable, it can be input to the network directly



$$x =$$

59
42
63
93
102
88
86
58
...

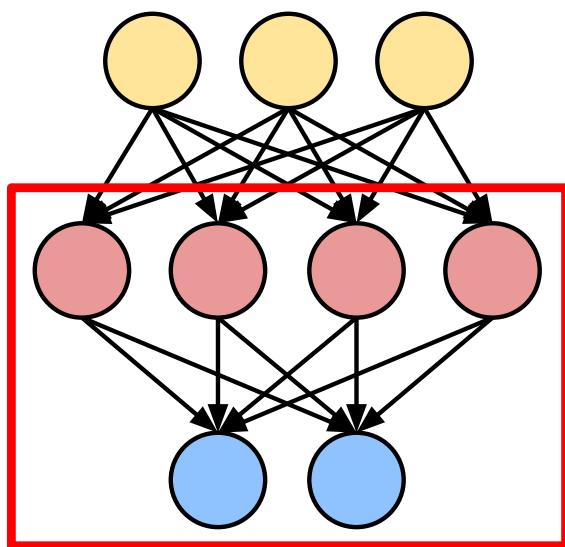


Last layer activation function

For classification, each neuron in the output represents one output class

You want the output to represent the probability that the input belongs to each class

What constraints must the output satisfy?



$$P(y = \text{"cat"})$$
$$P(y = \text{"dog"})$$

Last layer activation function

What constraints must the output satisfy?

- Probabilities are positive
- Probabilities sum to 1

$$\begin{aligned} P(y = \text{"cat"}) &\geq 0 \\ P(y = \text{"dog"}) &\geq 0 \end{aligned} \quad P(y = \text{"cat"}) + P(y = \text{"dog"}) = 1$$

Softmax activation function

$$\text{softmax}(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^n e^{z_j}}$$

Ensures positivity
Normalizes sum to 1

$$\mathbf{z} = \begin{bmatrix} 1.1 \\ 0.4 \\ -8.2 \\ -0.9 \\ 3.7 \end{bmatrix}$$
$$e^{\mathbf{z}} = \begin{bmatrix} e^{1.1} \\ e^{0.4} \\ e^{-8.2} \\ e^{-0.9} \\ e^{3.7} \end{bmatrix}$$
$$\frac{e^{\mathbf{z}}}{\sum_{j=1}^n e^{z_j}} = \begin{bmatrix} e^{1.1} / \sum_j e^{z_j} \\ e^{0.4} / \sum_j e^{z_j} \\ e^{-8.2} / \sum_j e^{z_j} \\ e^{-0.9} / \sum_j e^{z_j} \\ e^{3.7} / \sum_j e^{z_j} \end{bmatrix} = \begin{bmatrix} 0.07 \\ 0.03 \\ 0.00 \\ 0.01 \\ 0.89 \end{bmatrix}$$

Up Next:

Training neural networks

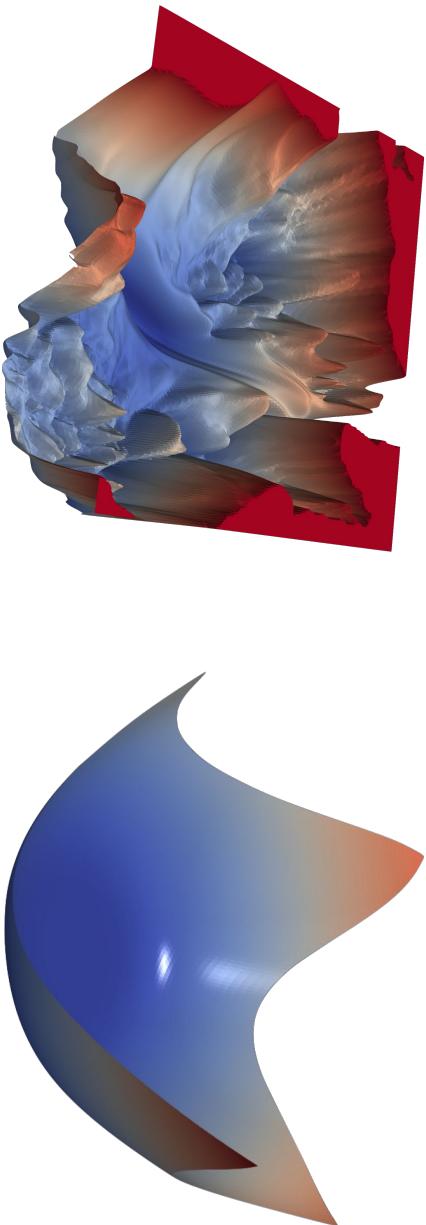
ICME Summer Workshops 2021

Introduction to Deep Learning

Session 2: 9:45–11:00 AM

Instructor: Sherrie Wang

icme-workshops.github.io/deep-learning



Li et al. "Visualizing the Loss Landscape of Neural Nets" (2018)

Workshop Schedule

Day 1

Session 1 (Today 8:00–9:30 AM)

- Introduction
- Examples of deep learning
- Math review
- Neural network basics

Day 2

Session 3 (Tomorrow 8:00–9:30 AM)

- Overfitting and underfitting
- Convolutional neural networks
- Recurrent neural networks

Session 4 (Tomorrow 9:45–11:00 AM)

- Other architectures
- Deep learning libraries
- Walkthrough of an example
- Failures of deep learning
- Walkthrough of an example

Poll:

Weight matrix

Go to tinyurl.com/dlworkshop2021

Suppose we have a neural network with 1 fully connected hidden layer. The input is a 32-dimensional vector. The output is a probability distribution across 8 classes. The weight matrix of the hidden layer should have dimension...

A. 32x1

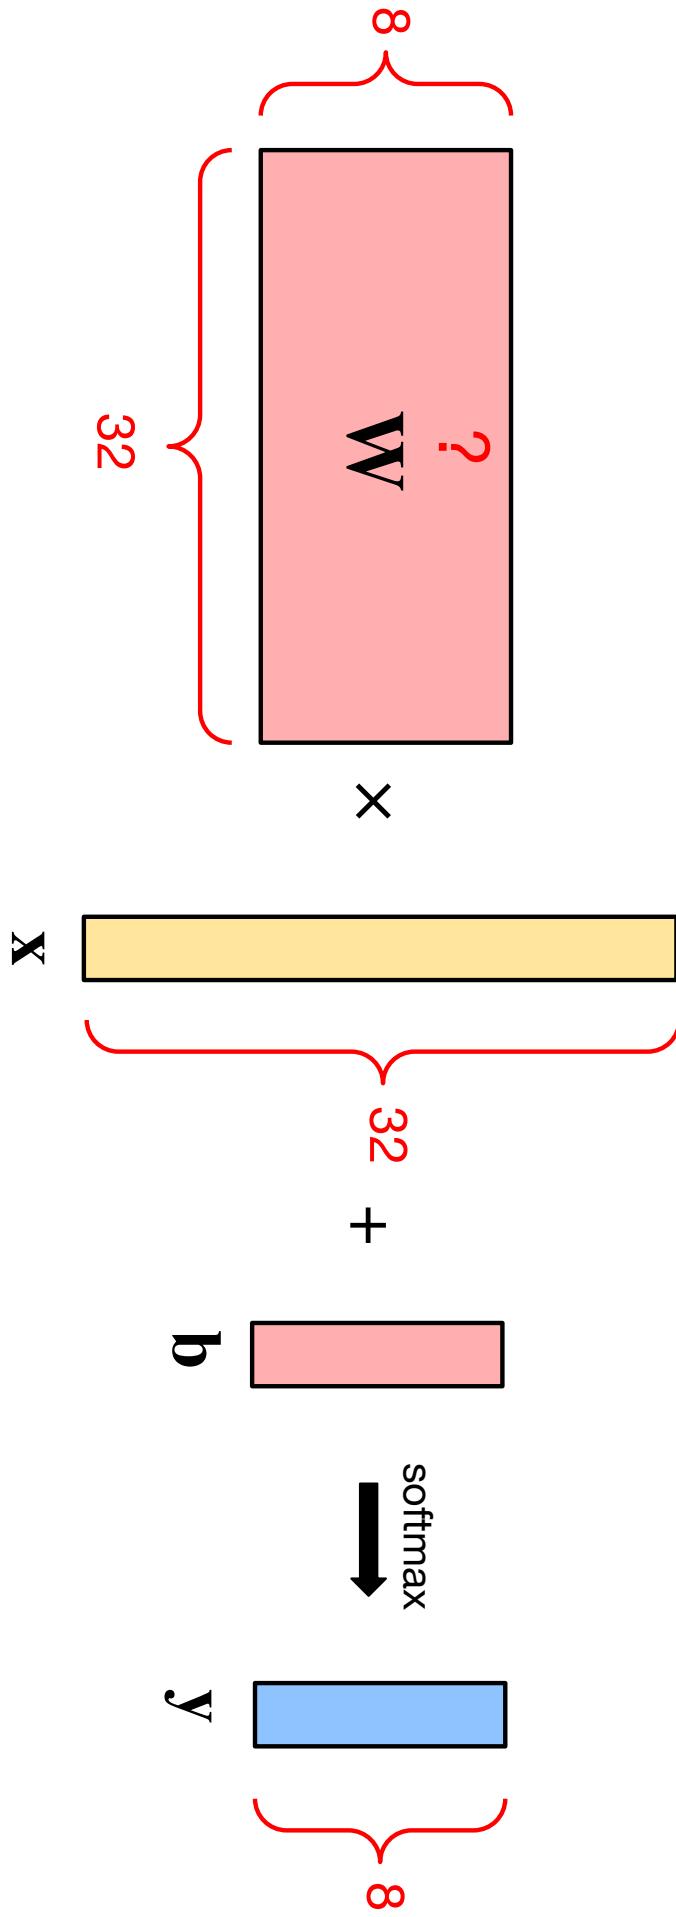
B. 8x32

C. 32x8

D. 32x32

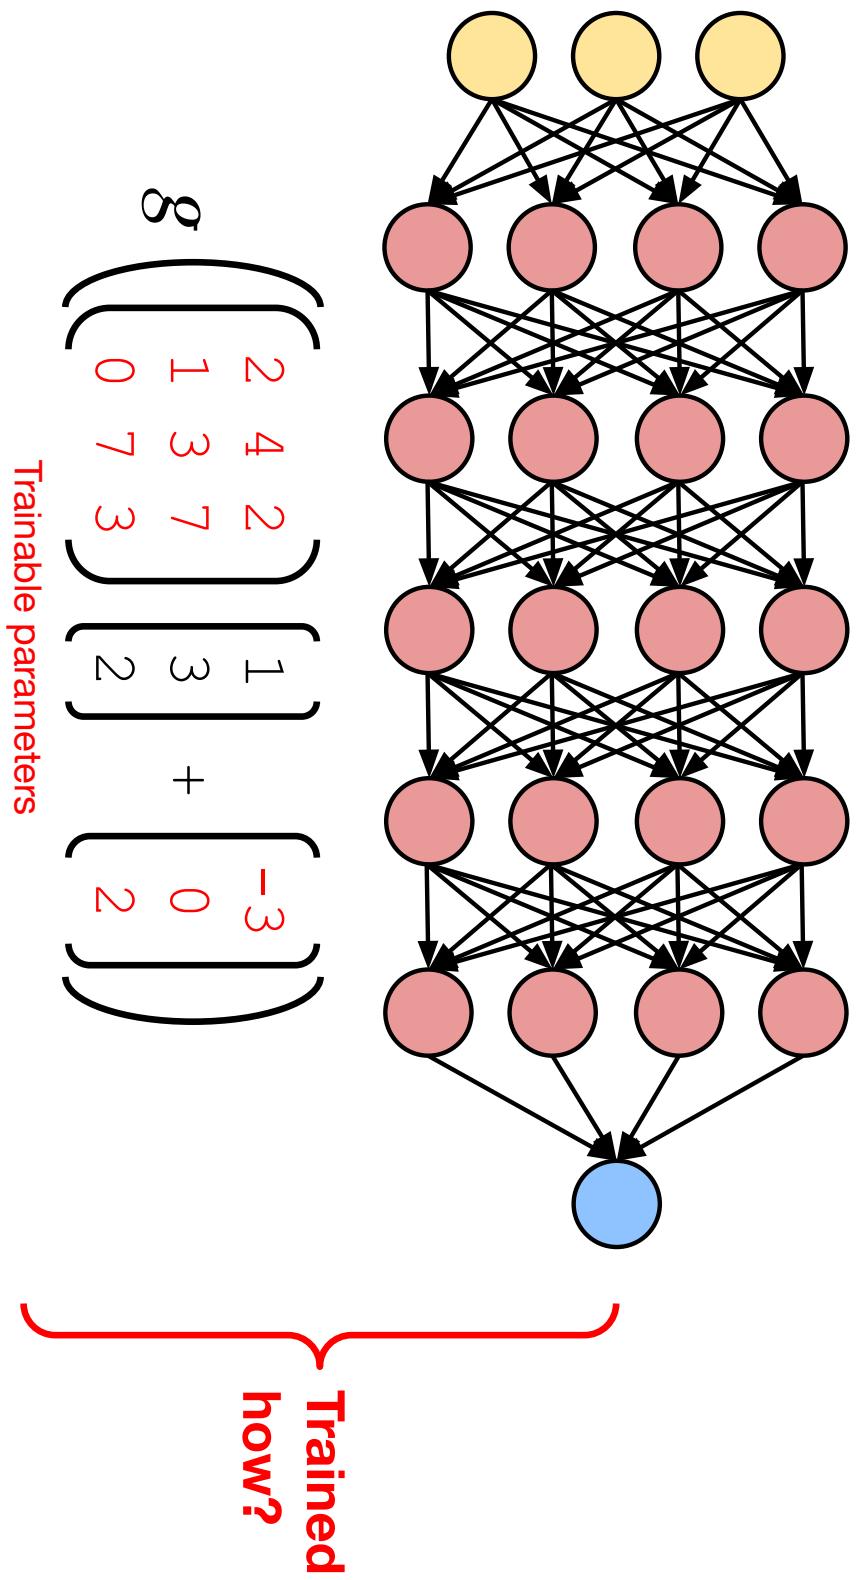
Poll:

Weight matrix



Loss Functions

Training neural networks

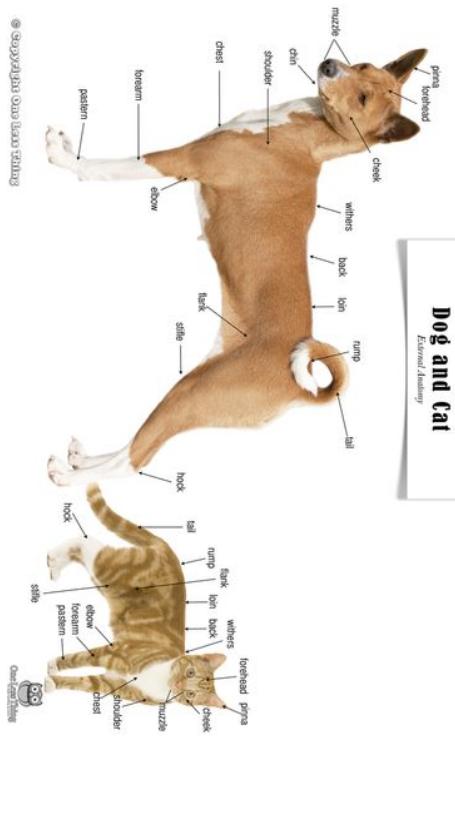


Training neural networks

Machine learning: builds a model based on “training data” in order to make predictions without being explicitly programmed to do so

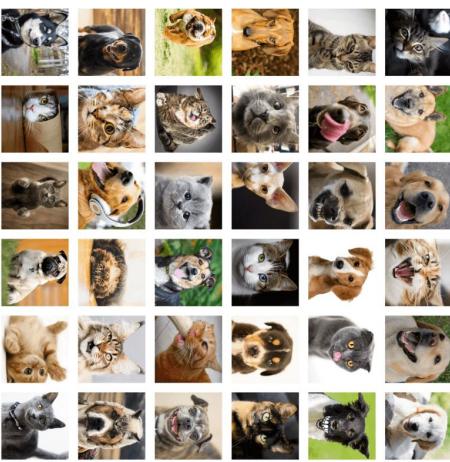
Not machine learning

Human says: “dogs are like this, and cats are like this...”



Machine learning

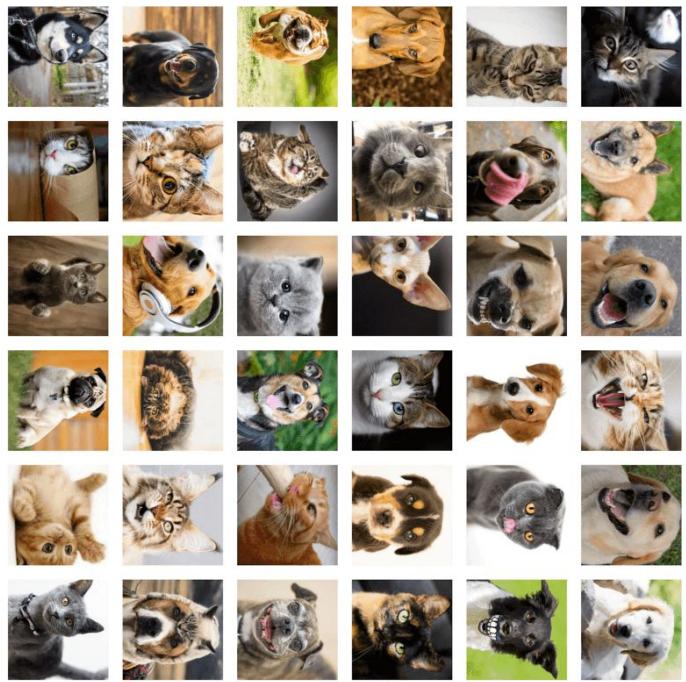
“Here’s a bunch of pictures of cats and dogs. Figure out how to sort them.”



<https://mc.ai/deep-learning-vs-machine-learning-a-simple-explanation/>

Training neural networks

Dataset



training

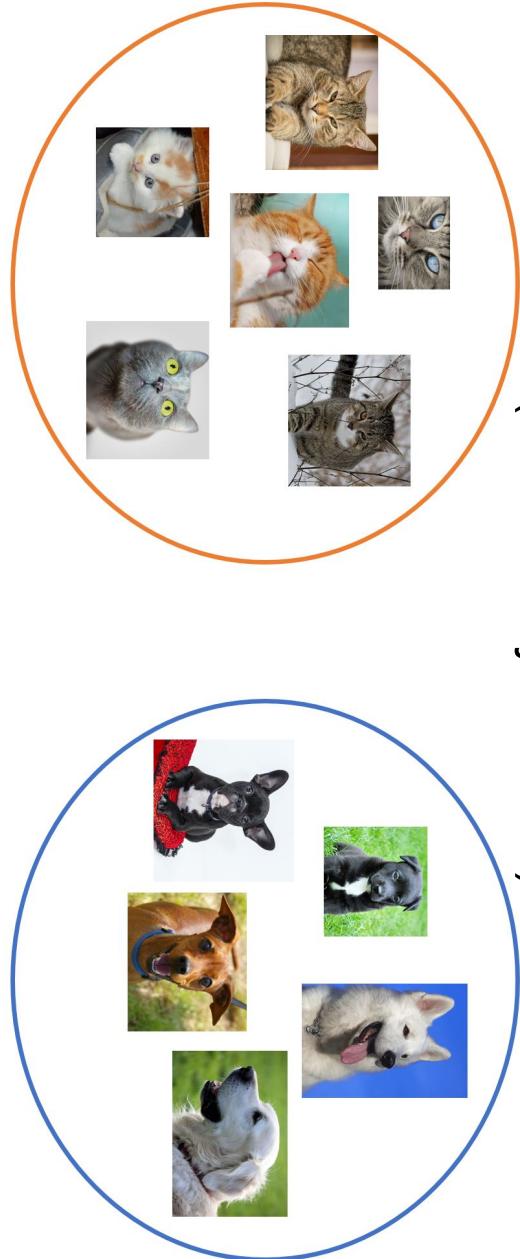
$$\begin{bmatrix} 2 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} 4 \\ 3 \\ 7 \end{bmatrix} \begin{bmatrix} 2 \\ 7 \\ 3 \end{bmatrix} \begin{bmatrix} -3 \\ 0 \\ 2 \end{bmatrix}$$

Weights & biases

Supervised learning

Dataset has labels
(Created by humans)

Cat
Dog



Supervised learning

MNIST Dataset (Handwritten digits)

0 0

1 1

2 2

3 3

4 4

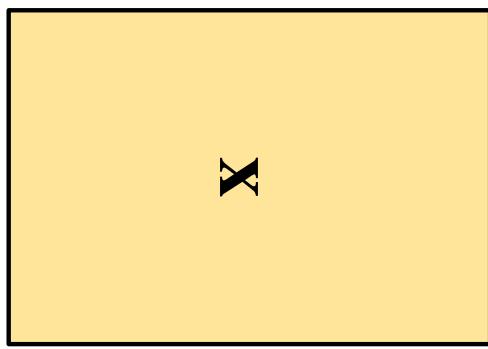
5 5

6 6

7 7

8 8

60,000 x 784



60,000 x 10



One-hot
encoding

Each image is 28 x 28 pixels

The goal of training

Intuitively, we want to tweak the weights of the neural network until the network predicts the correct output for most of the examples in the training set.

Neural network #1

$$\text{softmax} \left(\begin{bmatrix} 2 & 1 \\ -1 & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} \right) = \begin{bmatrix} 0.98 \\ 0.02 \end{bmatrix}$$

Weights Input Prediction

Neural network #2

$$\text{softmax} \left(\begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} \right) = \begin{bmatrix} 0.12 \\ 0.88 \end{bmatrix}$$

Weights Input Prediction

True label $\mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$ We prefer network #1 to #2

Optimization formulation

We will frame this as an optimization problem:

$$\max_{\mathbf{W}_i, \mathbf{b}_i} \text{accuracy}(\text{DNN}(\mathbf{W}_i, \mathbf{b}_i))$$

Accuracy is the fraction of training examples that the model predicts correctly

Optimization formulation

Accuracy is not a continuous function, so it is hard to directly optimize for maximum accuracy.

So we use **loss functions**.

Think of losses as continuous proxies to accuracy, so that we can formulate it as an optimization problem and solve it.

Loss functions

$$\begin{aligned} \text{True label } y &= \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} & \text{Predicted probabilities} &= \begin{bmatrix} 0 \cdot 2 \\ 0 \cdot 7 \\ 0 \cdot 1 \end{bmatrix} \end{aligned}$$

In English, we want to **reward the neural network** for predicting class 2 with higher probability than classes 1 and 3, while incentivizing it to become even better at making this distinction

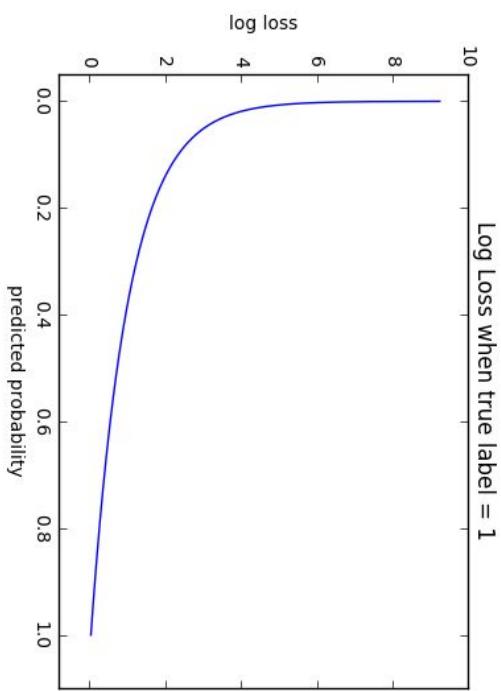
Cross entropy loss

$$\text{Predicted probabilities} = \begin{bmatrix} 0.2 \\ 0.7 \\ 0.1 \end{bmatrix} - \log(0.7) = 0.35$$

$$\text{Predicted probabilities} = \begin{bmatrix} 0.6 \\ 0.1 \\ 0.3 \end{bmatrix} - \log(0.1) = 2.3$$

Cross entropy loss

$$\text{CE}(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_i y_i \log \hat{y}_i$$



Cost function

Cost = Average loss over
training examples

$$\min_{\mathbf{W}_i, \mathbf{b}_i} \text{cost}(\text{DNN}(\mathbf{W}_i, \mathbf{b}_i))$$

Often written

$$\min_{\mathbf{W}_i, \mathbf{b}_i} \mathcal{J}(\mathbf{W}_i, \mathbf{b}_i)$$

The objective function is now a continuous function of the weights and biases

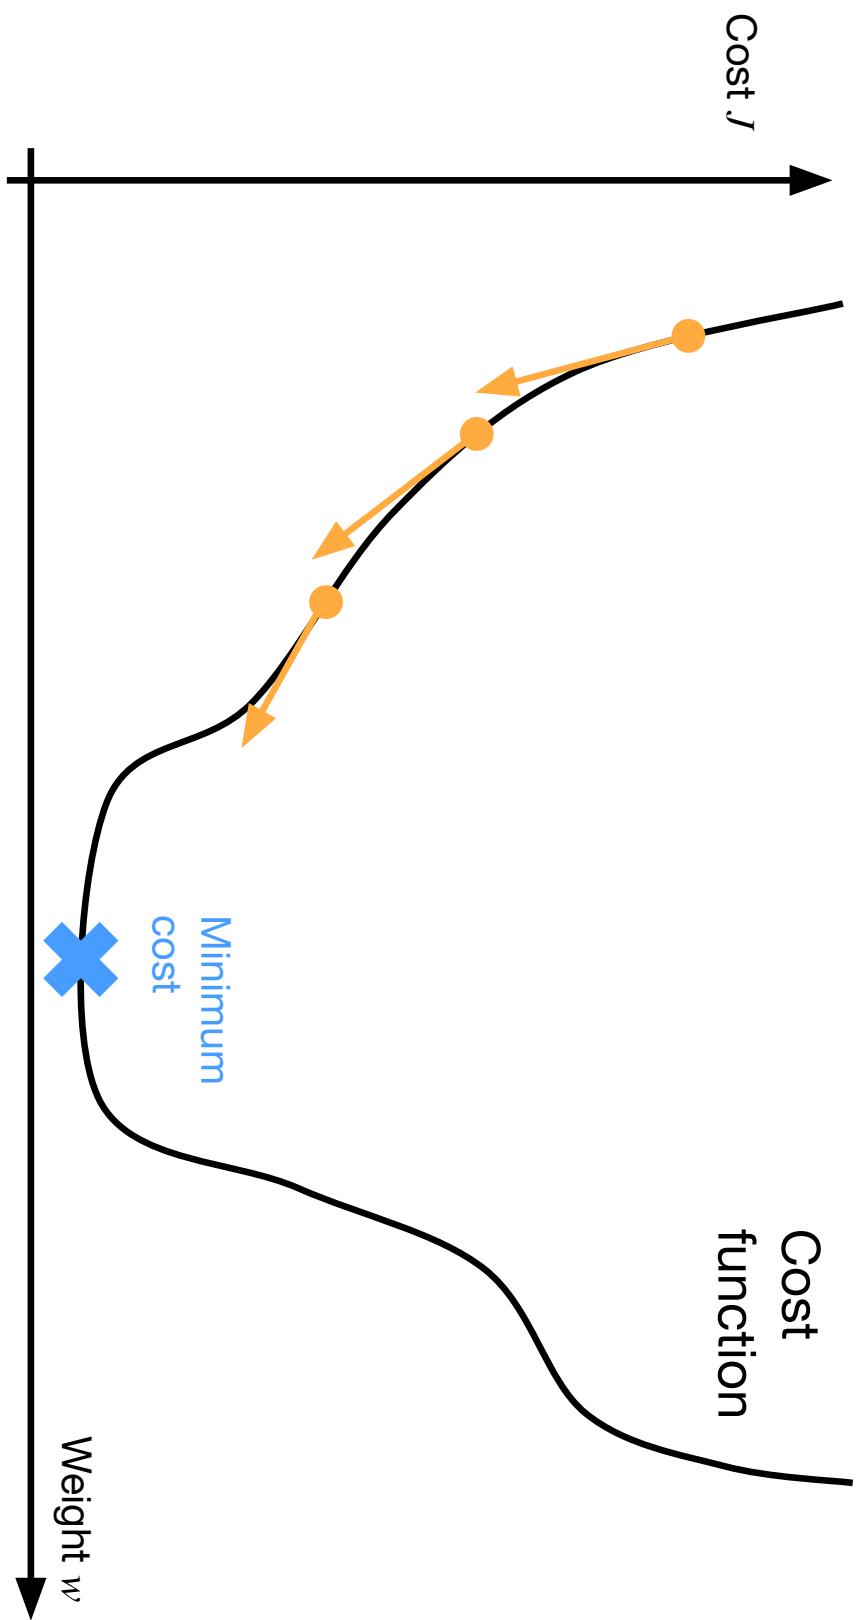
Gradient Descent

How do we solve the optimization problem?

The purpose of reformulating the problem in terms of a cost function is so that we can use gradient-based methods

1. Start with random weights
2. Compute the gradient with respect to the cost function
3. Update the weights so that the cost function decreases
4. Repeat steps 2-3

Gradient descent

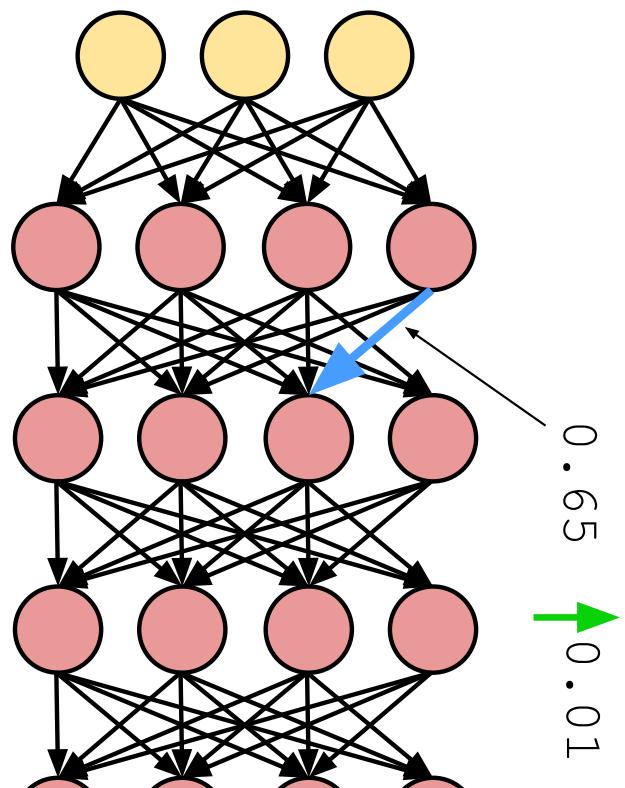


What are gradients?

1. It is a vector whose elements are the partial derivatives associated with every weight and bias in the neural network
2. The partial derivative on a weight = how many units the cost function will change by if the weight is changed by one unit
3. This assumes that the unit of change is small (technically infinitesimally small)

$$\nabla f(x, y, z) = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \\ \frac{\partial f}{\partial z} \end{bmatrix}$$

What are gradients?



$$J = 4.50$$
$$\downarrow 0.05$$

$$\text{Gradient} = \frac{\partial J}{\partial w_i} = \frac{-0.05}{0.01} = -5$$

What are gradients?

Suppose you have a function $F(x, y, z)$

Then we have partial derivatives F_x , F_y , and F_z

$$F(x + \Delta x, y, z) \approx F(x, y, z) + F_x \Delta x$$

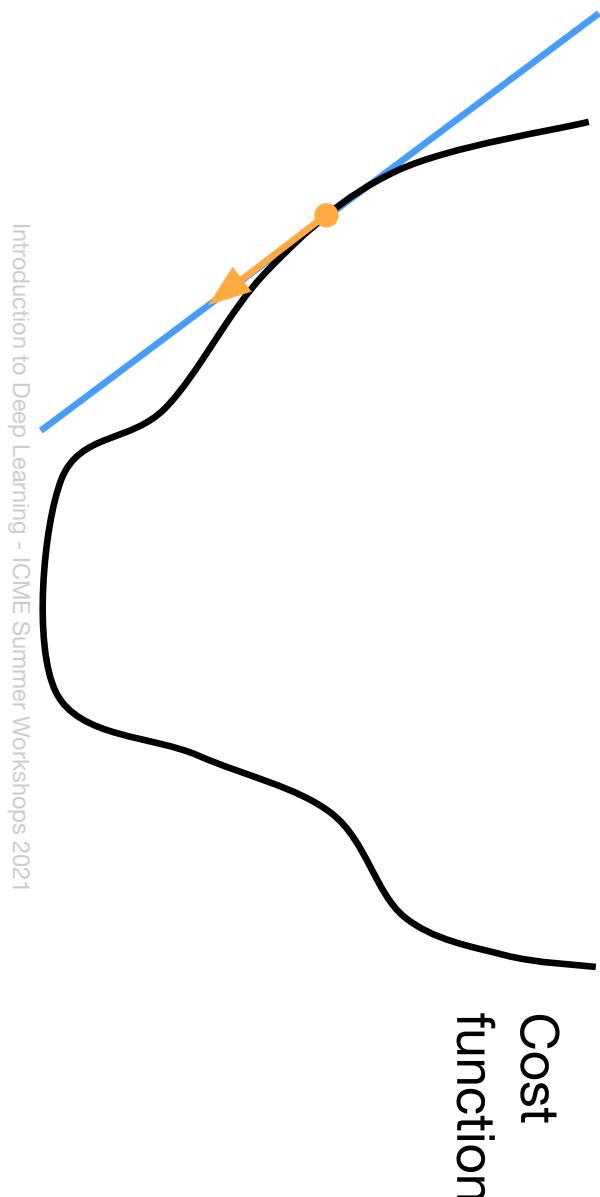
$$F(x, y + \Delta y, z) \approx F(x, y, z) + F_y \Delta y$$

$$F(x, y, z + \Delta z) \approx F(x, y, z) + F_z \Delta z$$

Linear approximation to the cost function

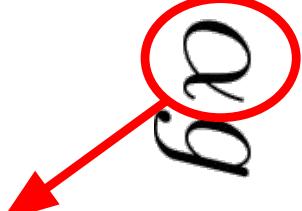
We can build a linear approximation to the function $F(x,y,z)$ using the gradients

$$F(x+\Delta x, y+\Delta y, z+\Delta z) \approx F(x, y, z) + F_x \Delta x + F_y \Delta y + F_z \Delta z$$



Gradient descent

1. Go downhill according to the linear approximation of the cost function
2. But not by too much, as the approximation may not be accurate far away
3. This means updating each weight in proportion to the gradient:

$$w' = w - \alpha g$$


Learning
rate

Backpropagation

How to compute gradients

- Numerically approximating gradients is too slow and computationally expensive
- The neural network is a well-defined mathematical function, so we should be able to differentiate it and calculate the derivatives using the **Chain Rule**
- This also too tedious, so we use an algorithm called **backpropagation**



$$\begin{aligned} & \frac{d}{dx} \sqrt{\sqrt{(x-1)(x+2)} + 1} \\ &= \frac{d}{dx} \left\{ [(x-1)(x+2)]^{\frac{1}{2}} + 1 \right\}^{\frac{1}{2}} \\ &= \frac{1}{2} \left\{ [(x-1)(x+2)]^{\frac{1}{2}} + 1 \right\}^{\frac{1}{2}-1} \frac{d}{dx} \left\{ [(x-1)(x+2)]^{\frac{1}{2}} + 1 \right\} \\ &= \frac{1}{2} \left\{ [(x-1)(x+2)]^{\frac{1}{2}} + 1 \right\}^{\frac{1}{2}-1} \frac{1}{2} \left\{ [(x-1)(x+2)]^{\frac{1}{2}-1} \frac{d}{dx} [(x-1)(x+2)] \right\} \\ \dots \text{like this?} &= \frac{1}{2} \left\{ [(x-1)(x+2)]^{\frac{1}{2}} + 1 \right\}^{\frac{1}{2}-1} \frac{1}{2} \left\{ [(x-1)(x+2)]^{\frac{1}{2}-1} \left\{ \frac{d}{dx} (x-1) [x+2] + (x-1) \left[\frac{d}{dx} (x+2) \right] \right\} \right\} \end{aligned}$$

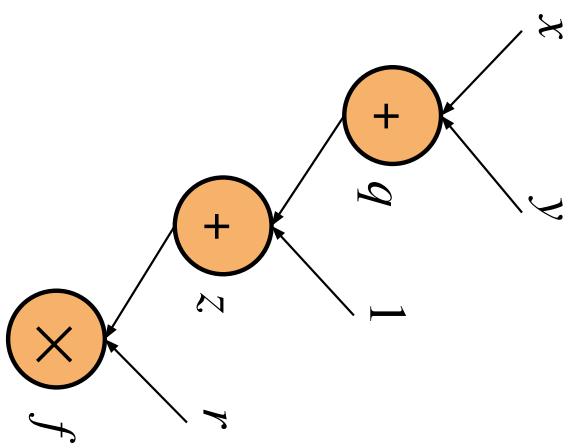
Computational graph

A way to organize mathematical expressions

Mathematical expressions

$$\begin{aligned}x + y &= q \\q + 1 &= z \\z \times r &= f\end{aligned}$$

Computational graph

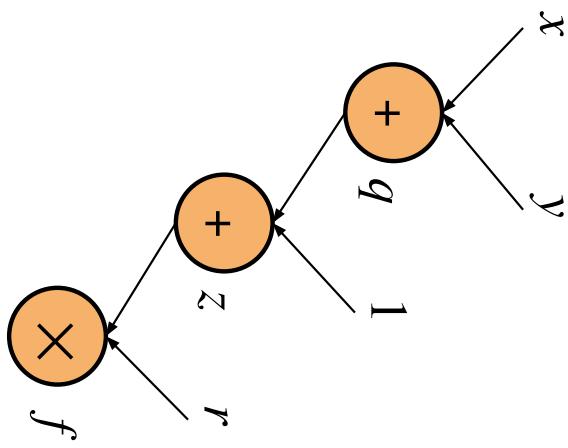


Each circle
denotes an
operation

Gradients on computational graphs

Suppose we want to know $\frac{\partial f}{\partial x}$.

We can find this by chaining together derivatives at each operation.

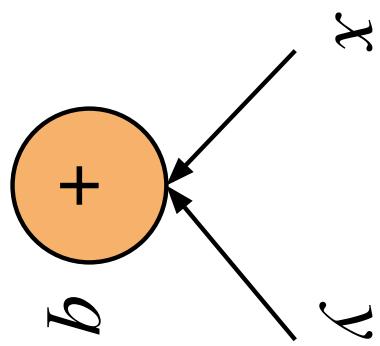


Gradients on computational graphs

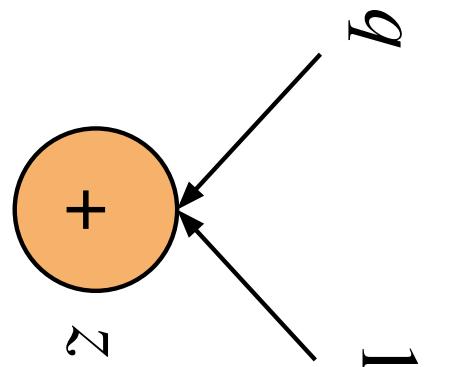
Let's take one piece of the computational graph at a time.

$$x + y = q$$

$$\frac{\partial q}{\partial x} = 1 \quad \frac{\partial q}{\partial y} = 1$$



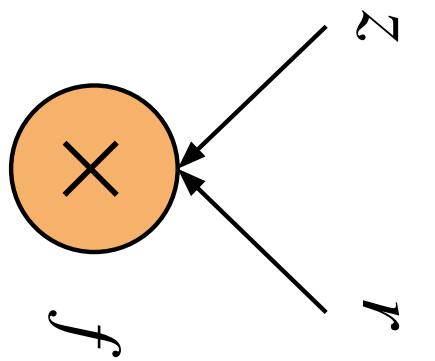
Gradients on computational graphs

$$q + 1 = z$$
$$\frac{dz}{dq} = 1$$


**Chain
Rule**

$$\frac{\partial z}{\partial x} = \frac{dz}{dq} \frac{\partial q}{\partial x} = 1$$

Gradients on computational graphs



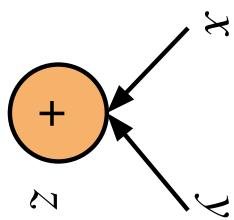
$$z \times r = f$$

$$\frac{\partial f}{\partial z} = r \quad \frac{\partial f}{\partial r} = z$$

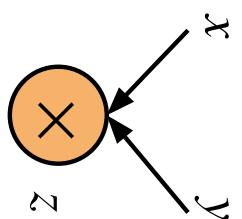
Chain Rule $\frac{\partial f}{\partial x} = \frac{\partial f}{\partial z} \frac{dz}{dq} \frac{\partial q}{\partial x} = r$

In backprop, gradients are highly local

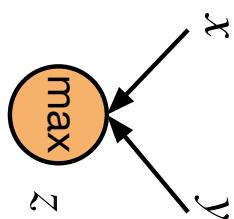
Three common operations in neural networks:



$$\frac{\partial f}{\partial x} = 1 \quad \frac{\partial f}{\partial y} = 1$$



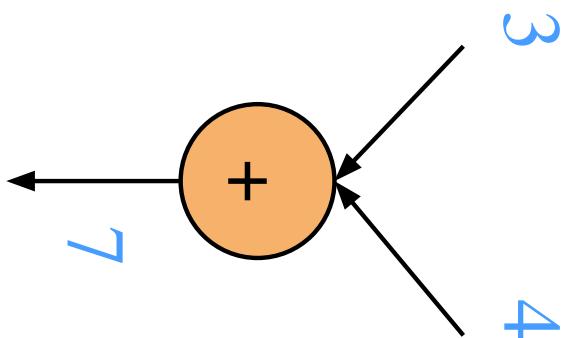
$$\frac{\partial f}{\partial x} = y \quad \frac{\partial f}{\partial y} = x$$



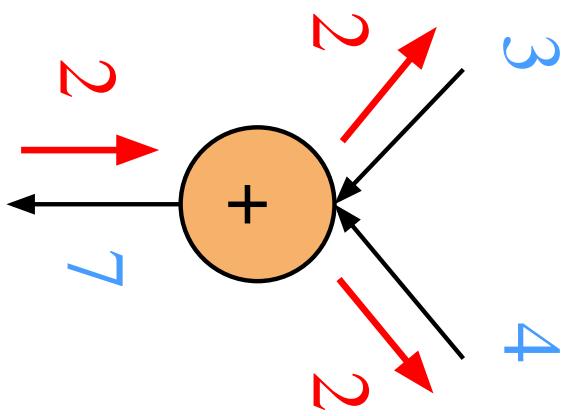
$$\frac{\partial f}{\partial x} = \mathbf{1}\{x \geq y\} \quad \frac{\partial f}{\partial y} = \mathbf{1}\{y \geq x\}$$

A small example with numbers

Forward pass



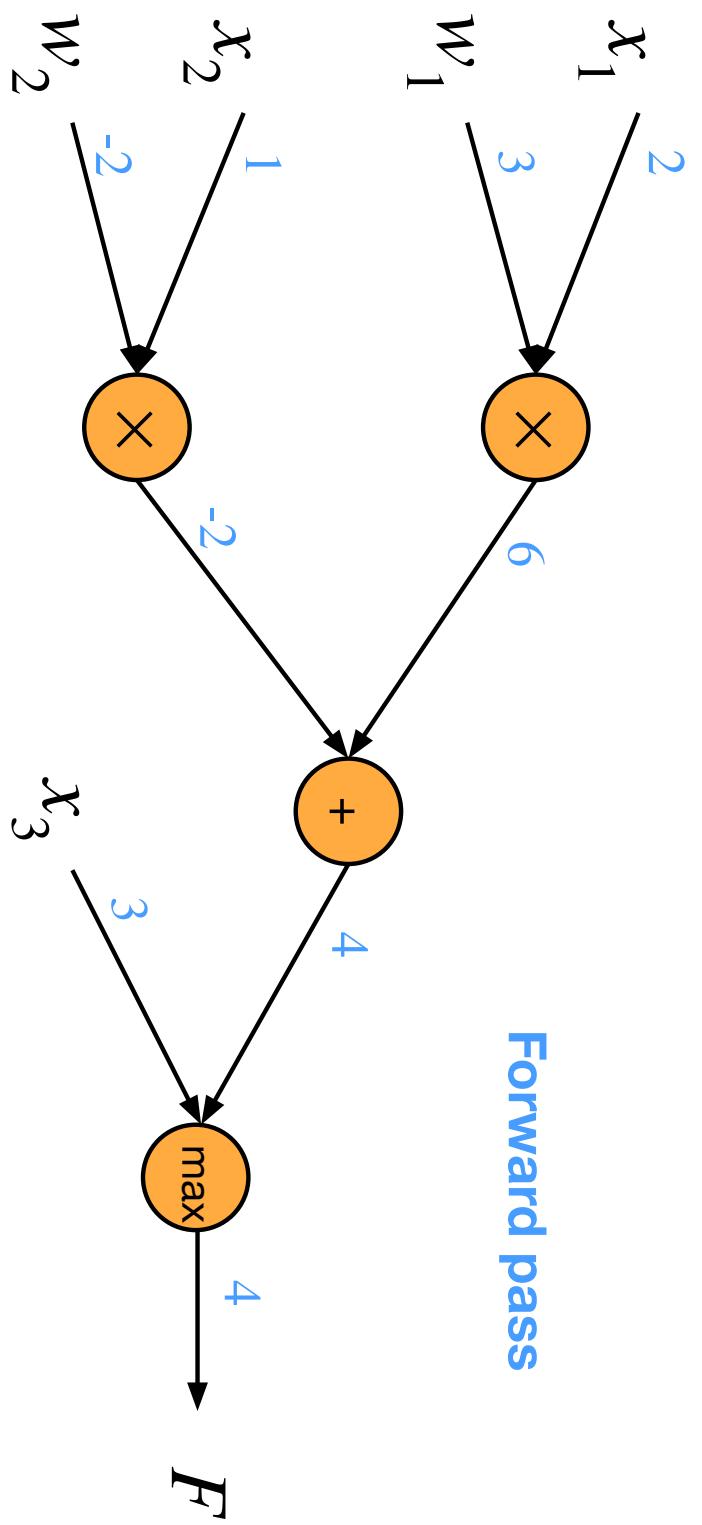
Backward pass



A larger example with numbers

$$F(x_1, x_2, x_3) = \max(w_1x_1 + w_2x_2, x_3)$$

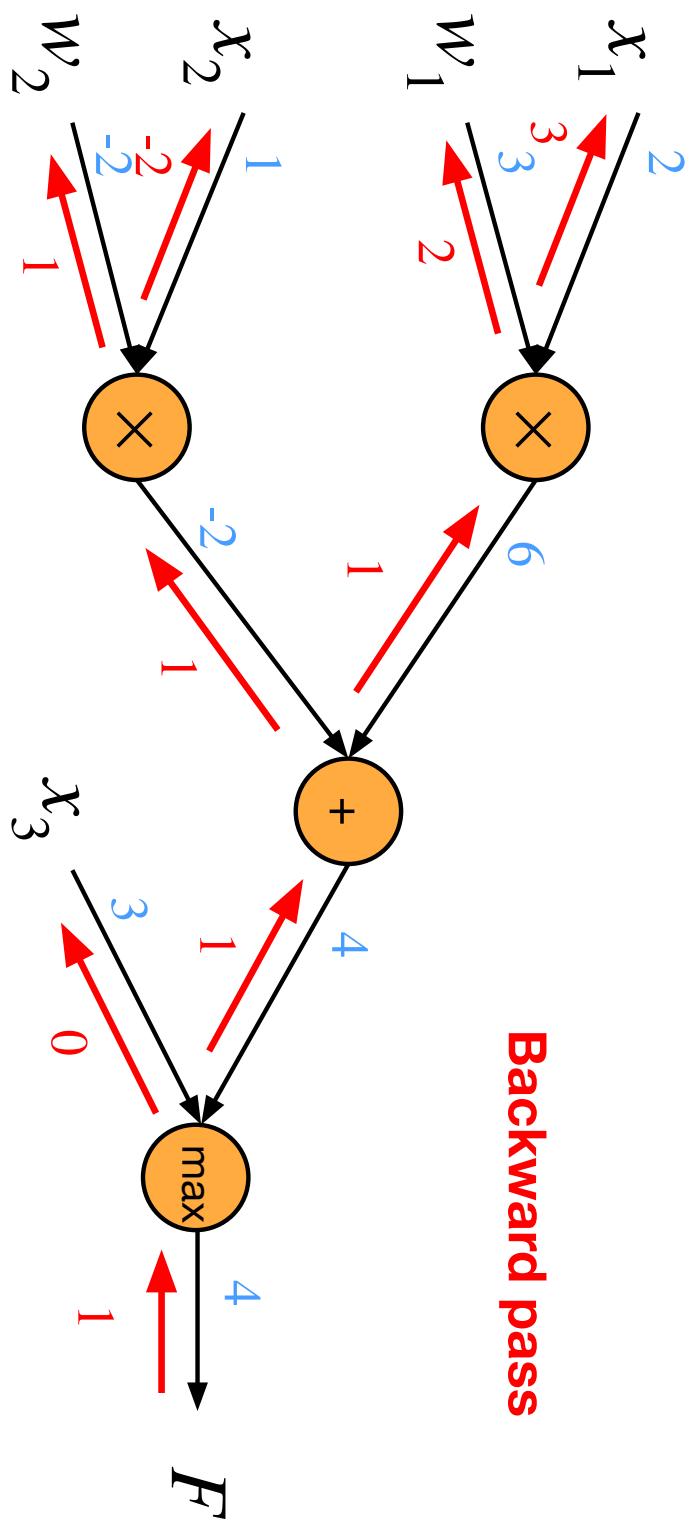
Forward pass



A larger example with numbers

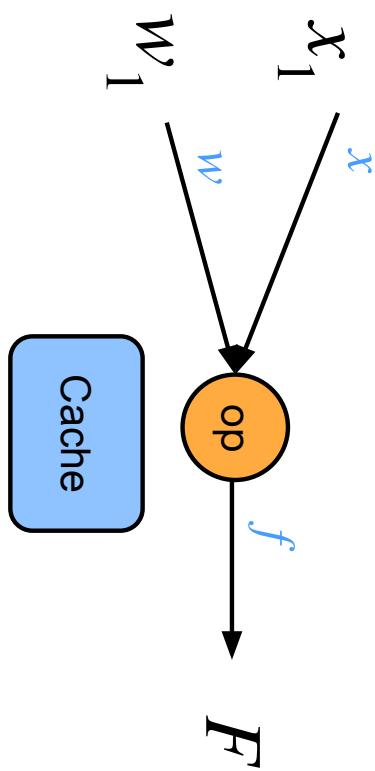
$$F(x_1, x_2, x_3) = \max(w_1x_1 + w_2x_2, x_3)$$

Backward pass

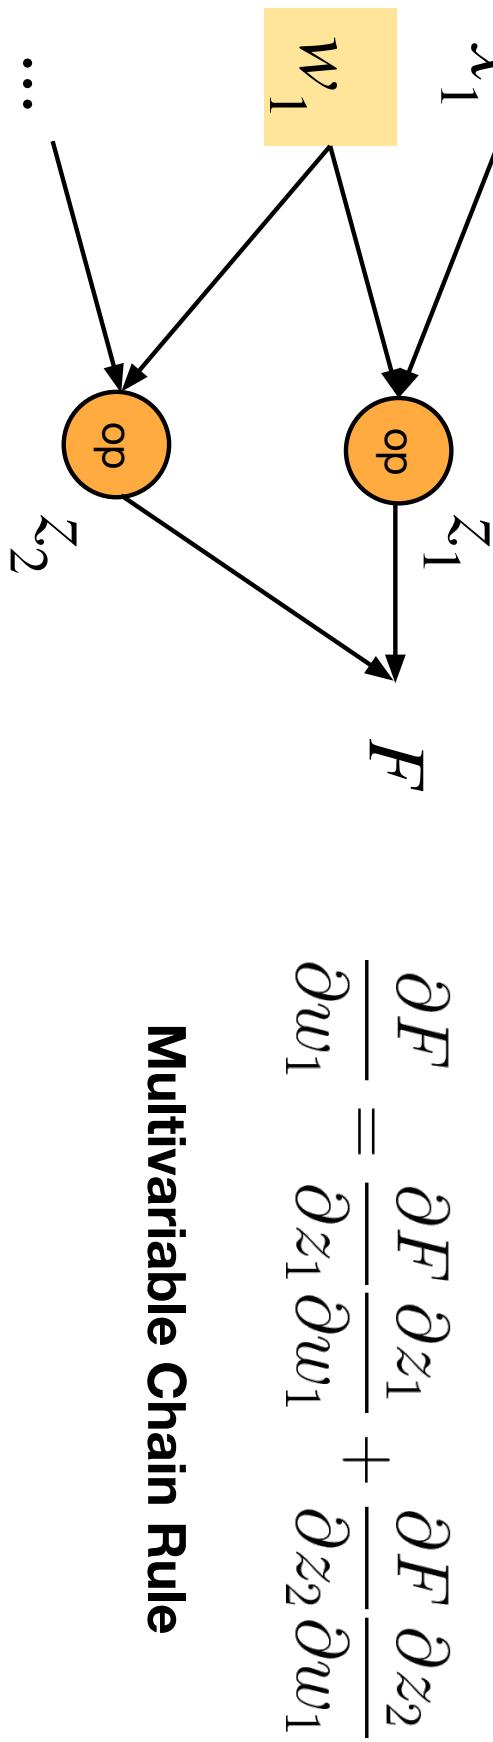


Note 1: Cache forward pass variables

During the backward pass, variable values from the forward pass are often needed

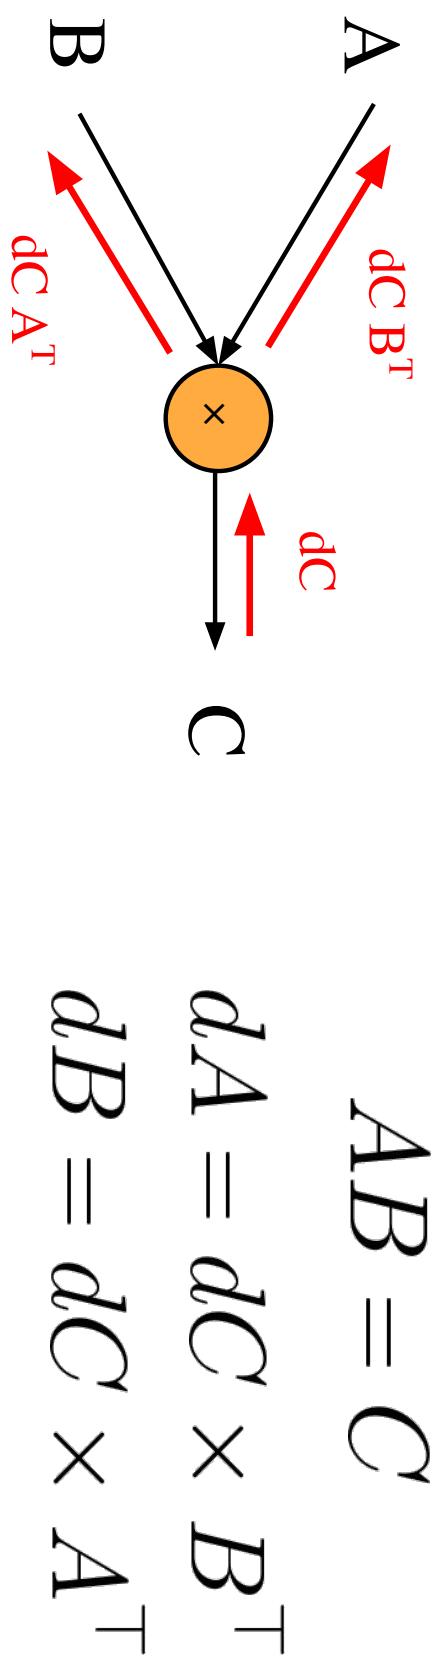


Note 2: Gradients add when a variable appears multiple times



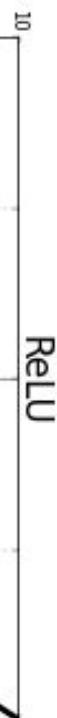
Multivariable Chain Rule

Note 3: Backprop through matrix multiplication



Check matrix dimensions to verify

Poll: What happens when the gradient passes through the ReLU function?



Go to tinyurl.com/dlworkshop2021

- A. It stays the same
- B. It is multiplied by a constant
- C. It becomes zero
- D. It either stays the same or is zeroed
- E. None of the above

Gradient descent

$$w' = w - \alpha g$$

The cost function is the average of the loss function over all training examples.

This means we have to forward and backward pass through all training examples before we take one step of gradient descent.

Stochastic gradient descent

In SGD, we estimate the gradient by evaluating it over only a small subset of the training examples.

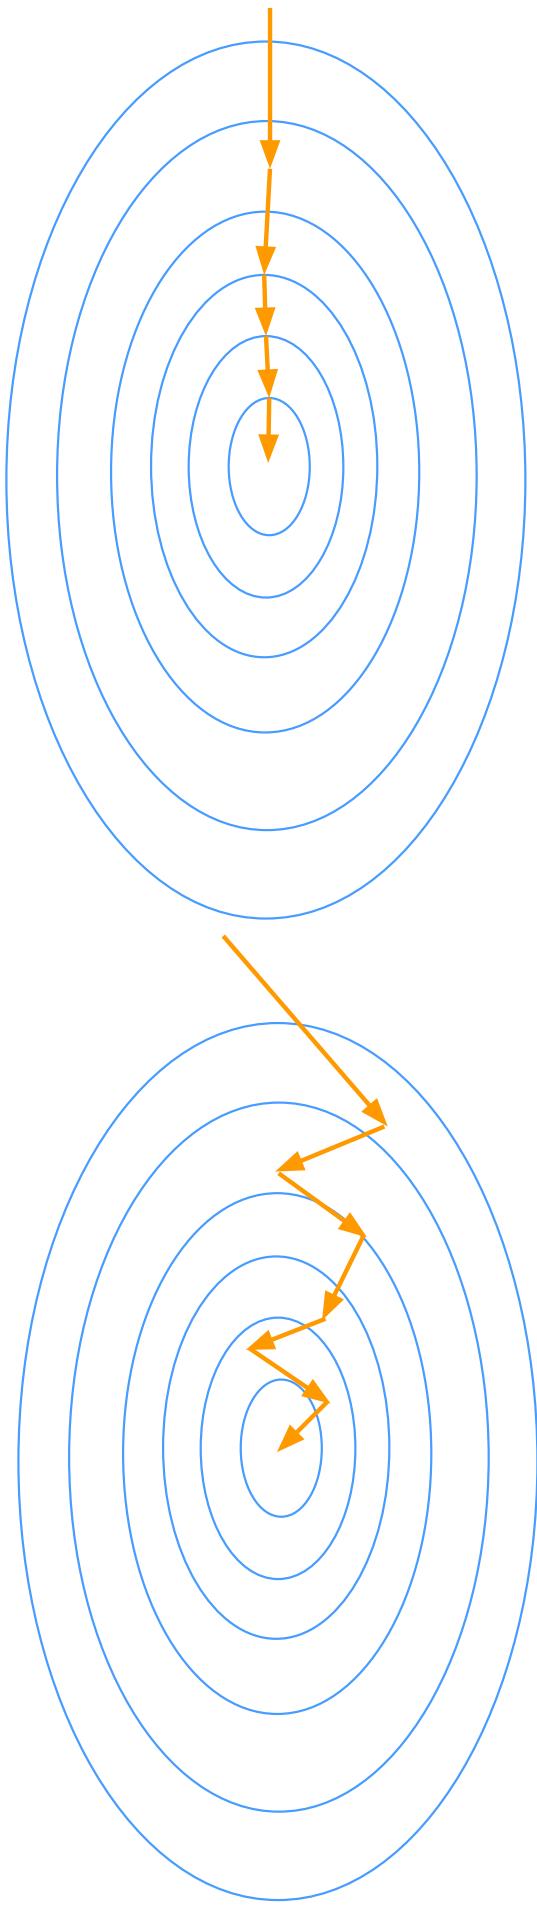
The size of this subsample is the **batch size**, which becomes another hyperparameter.

We use a different subsample for each gradient step, so as to go through all the training examples. Going through the entire training set once is called one **epoch** of training.

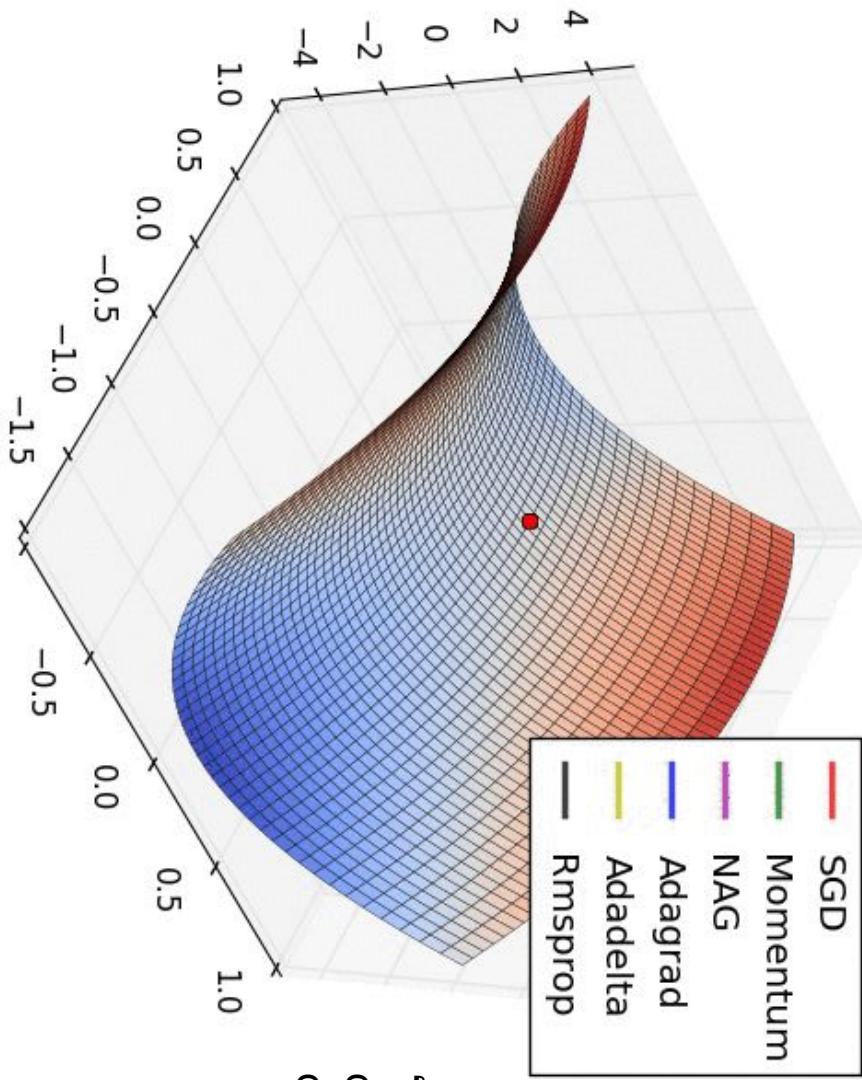
Stochastic gradient descent

Gradient descent

Stochastic gradient descent



Other optimization algorithms



“Adam” is a very
commonly used
optimizer

When to stop training neural network?

Until convergence, when the loss no longer changes much.

In practice, often for a fixed number of epochs.

What data do we train on?

Training, validation, and test sets

To evaluate the performance of the model, set aside a portion of the dataset.

Training set: used to train the weights of the neural network.

Validation set: used to tune the network hyperparameters and perform architecture search.

Test set: reserved until the end, used only once to evaluate the final weights and hyperparameters of the model.



<https://towardsdatascience.com/train-test-split-and-cross-validation-in-python-80b61beca4b6>

Poll:

Data splits

Go to tinyurl.com/dlworkshop2021

Suppose I want to try training my neural network with different learning rates $\{0.1, 0.01, 0.001, 0.0001\}$. Which split should I use to evaluate which learning rate is best?

- A. Train
- B. Validation
- C. Test

Training loop

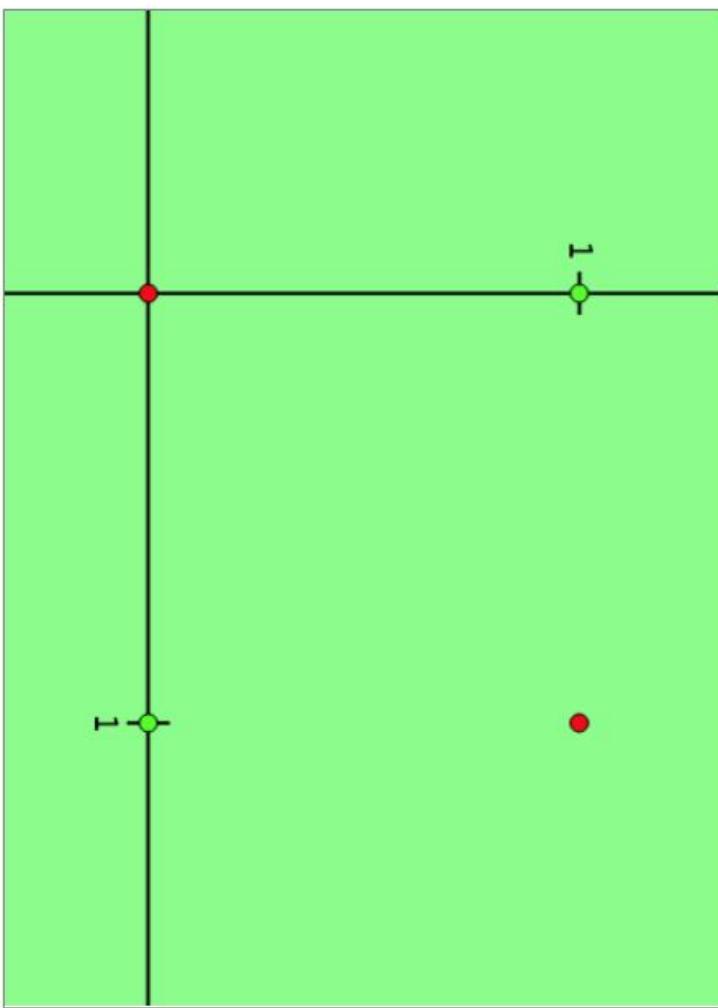
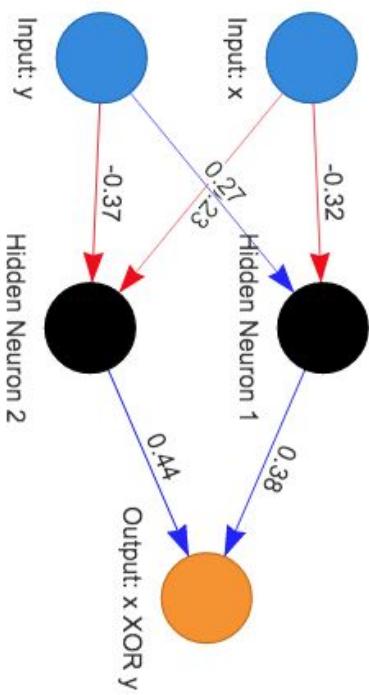
1. Sample a batch of data
2. Forward pass through all examples in the batch
3. Compute loss
4. Backpropagate through all examples
5. Compute final gradient on batch and update weights
6. Repeat

Workflow

1. Split your data into train/val/test
2. Choose architecture
3. Choose hyperparameters: learning rate, batch size, etc
4. Train neural network until convergence, evaluating on validation set each epoch
5. Check if training is overfitting or underfitting
6. Modify architecture and hyperparameters as needed, repeat

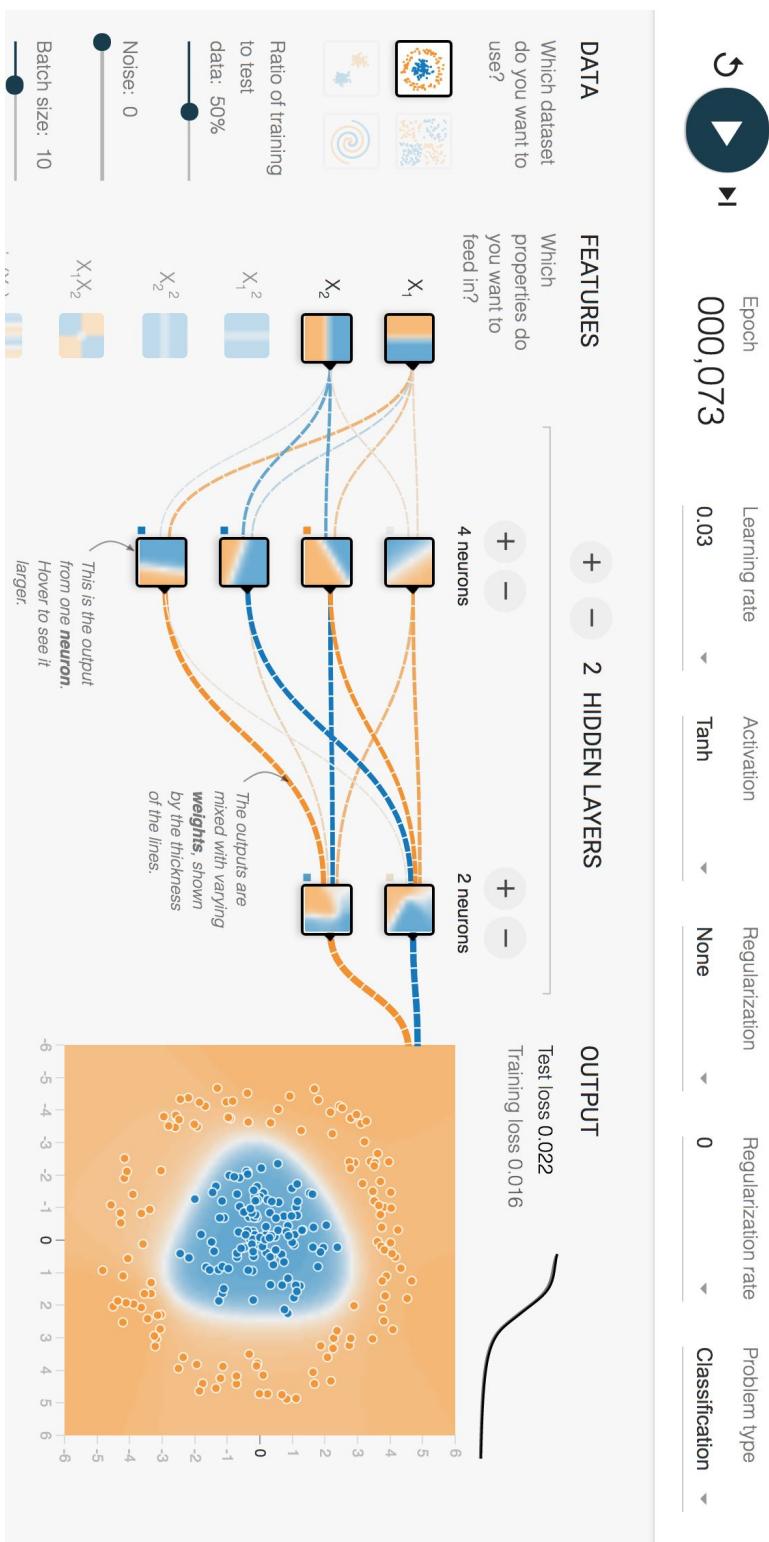
A complete example

<https://phiresky.github.io/neural-network-demo/>



TensorFlow playground

<https://playground.tensorflow.org>



That's it for today! See you tomorrow!

Day 1

Session 1 (Today 8:00–9:30 AM)

- Introduction
- Examples of deep learning
- Math review
- Neural network basics

Session 2 (Today 9:45–11:00 AM)

- Loss functions
- Gradient descent
- Backpropagation
- Walkthrough of an example

Day 2

Session 3 (Tomorrow 8:00–9:30 AM)

- Overfitting and underfitting
- Convolutional neural networks
- Recurrent neural networks

Session 4 (Tomorrow 9:45–11:00 AM)

- Other architectures
- Deep learning libraries
- Walkthrough of an example
- Failures of deep learning